

nf-core/taxprofiler: highly parallelised and flexible pipeline for metagenomic taxonomic classification and profiling

Sofia Stamouli¹, Moritz E. Beber², Tanja Normark³, Thomas A. Christensen II⁴, Lili Andersson-Li⁵, Maxime Borry⁶, Mahwash Jamy⁷, nf-core community⁸, James A. Fellows Yates⁹

¹Department of Microbiology, Tumor and Cell Biology, Karolinska Institutet

¹Department of Clinical Microbiology, Karolinska University Hospital

²Unseen Bio ApS

³Department of Microbiology, Tumor and Cell Biology, Karolinska Institutet

³Department of Clinical Microbiology, Karolinska University Hospital

⁴Veterinary Diagnostic Laboratory, Kansas State University College of Veterinary Medicine

⁵Department of Microbiology, Tumor and Cell Biology, Karolinska Institutet

⁵Department of Clinical Microbiology, Karolinska University Hospital

⁶Department of Archaeogenetics, Max Planck Institute for Evolutionary Anthropology

⁷Department of Microbiology, Tumor and Cell Biology, Karolinska Institutet

⁷Department of Clinical Microbiology, Karolinska University Hospital

⁸

⁹Department of Archaeogenetics, Max Planck Institute for Evolutionary Anthropology

⁹Department of Paleobiotechnology, Leibniz Institute for Natural Product Research and Infection Biology Hans Knöll Institute

1 Abstract

Metagenomic classification tackles the problem of characterising the taxonomic source of all DNA sequencing reads in a sample. A common approach to address the differences and biases between the many different taxonomic classification tools is to run metagenomic data through multiple classification tools and databases. This, however, is a very time-consuming task when performed manually - particularly when combined with the appropriate preprocessing of sequencing reads before the classification.

Here we present nf-core/taxprofiler, a highly parallelised read-processing and taxonomic classification pipeline. It is designed for the automated and simultaneous classification and/or profiling of both short- and long-read metagenomic sequencing libraries against a 11 taxonomic classifiers and profilers as well as databases within a single pipeline run. Implemented in Nextflow and as part of the nf-core initiative, the pipeline benefits from high levels of scalability and portability, accommodating from

36 small to extremely large projects on a wide range of computing infrastructure. It has
37 been developed following best-practise software development practises and commu-
38 nity support to ensure longevity and adaptability of the pipeline, to help keep it up to
39 date with the field of metagenomics.

40 2 Introduction

41 Whole-genome, metagenomic sequencing offers strong benefits to the taxonomic clas-
42 sification of DNA samples over targeted approaches (Eloe-Fadrosh et al. 2016; Florian
43 P. Breitwieser, Lu, and Salzberg 2019). While metabarcoding approaches targeting
44 the 16S rRNA or other marker genes are widely used due to low cost and large, di-
45 verse reference databases (Yilmaz et al. 2014; Lynch and Neufeld 2015), metagenomic
46 approaches have been gaining popularity with the increasingly lower costs of, for
47 example, shotgun sequencing. These metagenomic analyses with whole microbial
48 genome as references have been shown to provide a similar level of taxonomic res-
49 olution (Hillmann et al. 2018). However they also have the added benefit of having
50 greater reusability potential of the data, such as for whole genome and/or functional
51 classification (Sharpton 2014; Quince et al. 2017).

52 Taxonomic classifiers (sometimes referred to as taxonomic bidders) aim to identify
53 the original ‘taxonomic source’ of a given DNA sequence (Ye et al. 2019; Meyer et al.
54 2022; Govender and Eyre 2022). In metagenomics, this typically consists of comparing
55 millions of DNA reads (sequenced DNA molecules) against hundreds or thousands of
56 reference genomes either via sequence alignment or ‘k-mer matching’ (Sharpton 2014;
57 Sun et al. 2021). The reference genome with the most similar match to the read is then
58 considered the most likely original ‘source’ organism of that sequence. In this article
59 we will also refer to ‘taxonomic profilers’. We consider in these as classifiers that also
60 try to infer *species* or ‘biological’ abundance of the organism in the original sample, in
61 addition to the typical sequence abundance (Nayfach and Pollard 2016). We will use
62 classifiers and profilers interchangeably throughout the publication.

63 Having to identify the original source of the many DNA sequences out of the many ref-
64 erence genomes in a time and computationally efficient manner is a difficult problem.
65 In many cases biologists are not just interested as to which organism of each DNA
66 sequence comes from, but also in using this information to infer the original natu-
67 ral abundance of each organism of the given environment - something that is very
68 difficult due to the biases inherent to DNA extraction and sequencing. Therefore a
69 plethora of tools have been developed to address these challenges, all with their own
70 biases and specific contexts (Sczyrba et al. 2017; Meyer et al. 2022). Furthermore,
71 each tool often produces tool-specific output formats making it difficult to efficiently
72 cross compare results. Thus, no established ‘gold standard’ classifier tool or method
73 currently exists.

74 One solution to addressing the problem of choice among the range of different tools
75 is to run all of them in parallel, and cross compare the results. This can be useful both
76 for benchmarking studies (e.g. Sczyrba et al. 2017; Meyer et al. 2022), but also to

77 build consensus profiles whereby confidence of a particular taxonomic identification
78 can be increased when it is detected by multiple tools (McIntyre et al. 2017; Ye et al.
79 2019).

80 A second challenge in taxonomic classification (and arguably a larger one) is a ques-
81 tion of databases. As with tools, there is no one set ‘gold standard’ database. Different
82 questions and contexts require different databases, such as when a researcher wants
83 to search for both bacterial and viral species in samples, but as an extension of this,
84 taxonomic classifiers often will need different settings for each database. Further-
85 more, as genomic sequencing becomes cheaper and more efficient, the number of
86 publicly available reference genomes is rapidly increasing (Nasko et al. 2018). Conse-
87 quently, the size of reference databases of taxonomic classifiers is also growing, often
88 outpacing the computational capacity available to researchers. In fact, while this was
89 one of the main motivations behind classifiers such as Kraken2 (Wood, Lu, and Lang-
90 mead 2019), these algorithmic techniques are already becoming insufficient (Wright,
91 Comeau, and Langille 2023).

92 Finally, with the decrease of costs, the possibility for larger and larger metagenomic
93 sequencing datasets increases, leading to increasing sample sizes in studies. This is
94 exemplified by the doubling of the number of metagenomes on the European Bioin-
95 formatic Institute’s MGnify database within just two years (Mitchell et al. 2019).

96 Altogether this highlights the need for methods to efficiently profile many samples
97 using many tools. Manually setting up bioinformatic jobs for classification tasks for
98 each database and settings against different tools on traditional academic computing
99 infrastructure (e.g. high performance computing clusters or ‘HPC’ clusters) can be
100 very tedious. Additionally, particularly for very large sample sets, there is increas-
101 ing use of cloud platforms that have greater scalability than traditional HPCs. Being
102 able to reliably and reproducibly execute taxonomic classification tasks across infras-
103 tructure with minimal intervention would therefore be a boon for the metagenomics
104 field.

105 In recent years, workflow managers such as Nextflow (Di Tommaso et al. 2017) or
106 Snakemake (Mölder et al. 2021) have become highly popular in bioinformatics. These
107 frameworks provide for developers robust workflow execution with different HPC
108 scheduling tools and software provisioning systems, ensuring maximum portability
109 and efficient in different computational contexts. While a range of metagenomic
110 pipelines already exist (a non-exhaustive list being for example, StaG-mwc by
111 Boulund et al. 2023; MetaMeta by Piro, Matschkowski, and Renard 2017; TAMA by
112 Sim et al. 2020; UGENE by Rose et al. 2019; and Sunbeam by Clarke et al. 2019), few
113 leverage workflow managers to make multi-step workflows easier to use in HPC or
114 cloud infrastructure. Furthermore, often these pipelines aim to carry out multiple
115 different types of metagenomic analyses (e.g. also performing functional or assembly
116 analyses, such as Morais et al. 2022; Boulund et al. 2023) of which each step has
117 fewer options of tools and may execute functionality unwanted by the end user.

118 Here we present nf-core/taxprofiler (<https://nf-co.re/taxprofiler>), a pipeline designed
119 to allow users to efficiently and simultaneously taxonomically classify or profile

120 short- and long-read sequencing data. At the time of writing it supports 11 clas-
121 sifiers and an arbitrary number of databases per classifier in a single pipeline run.
122 nf-core/taxprofiler utilises Nextflow (Di Tommaso et al. 2017) to ensure efficiency,
123 portability, and scalability, and has been developed within the nf-core initiative of
124 Nextflow pipelines (Ewels et al. 2020) to ensure high quality coding practises and
125 user accessibility. It includes detailed documentation and a graphical-user-interface
126 (GUI) execution interface in addition to a standard command-line-interface (CLI).

127 3 Description

128 nf-core/taxprofiler aims to facilitate three main steps of a typical whole-genome,
129 metagenomic sequencing analysis workflow (Chiu and Miller 2019, Figure 1). A
130 longer description of the available functionality and motivations can be seen in the
131 [Supplementary Information](#).

132 In brief, nf-core/taxprofiler can accept short- (e.g. Illumina) and/or long-read
133 (e.g. Nanopore) FASTQ or FASTA files. These are supplied to the pipeline in the
134 form of a TSV file that includes basic sample and sequencing library metadata. The
135 pipeline can then be executed either via a standard Nextflow command-line-interface
136 execution or graphical-user-interface through either the open-source and free nf-core
137 launch page (<https://nf-core/launch>) or the commercial (with free-tier) Nextflow
138 tower (<https://tower.nf>) solution. Examples of the command-line execution and
139 nf-core launch GUI can be seen in the [Supplementary Information](#).

140 The pipeline can perform a range of metagenomics appropriate read preprocessing
141 steps, such as adapter removal, read merging, low-sequence complexity filtering, host-
142 or contamination removal, and/or per-sample run merging. All of these steps are
143 optional, and are aimed at removing possible sequencing artefacts that may result in
144 false positive taxonomic classification hits or improve classification efficiency. Most
145 of these steps also provide options of different tools to account for user preference.

146 After pre-processing, nf-core/taxprofiler can perform simultaneous profiling of pre-
147 processed reads with up to as many as 11 different taxonomic classifiers or profilers
148 (Table 1). Additionally on top of this, also simultaneously for each of the classifiers,
149 an arbitrary number of databases as supplied by the user. As of version 1.1.0, the
150 following classifiers and profilers are available: Kraken2 (Wood, Lu, and Langmead
151 2019), Bracken (Lu et al. 2017), KrakenUniq (F. P. Breitwieser, Baker, and Salzberg
152 2018), Centrifuge (Kim et al. 2016), MALT (Vågane et al. 2018), DIAMOND (Buchfink,
153 Reuter, and Drost 2021), Kaiju (Menzel, Ng, and Krogh 2016), MetaPhlAn (Blanco-
154 Míguez et al. 2023), mOTUs (Ruscheweyh et al. 2022), ganon (Piro et al. 2020), and
155 KMCP (Shen et al. 2023). Databases are also supplied via a input TSV file, which
156 also allows per-database custom classification parameters - meaning a given database
157 can be supplied multiple times each with different parameters or multiple different
158 databases per profiler. All classifiers with secondary steps to generate or convert to
159 additional output file formats are also included.

160 Post-processing of taxonomic profiles include standardisation and aggregation of pro-

161 files , i.e. merging of multiple profiles into a single multi-sample table for easier com-
 162 parison between profilers, with the tool TAXPASTA (Beber et al. 2023), and visualisa-
 163 tion of profiles with Krona (Ondov, Bergman, and Phillippy 2011) where supported.
 164 All relevant preprocessing statistics are displayed in an interactive and dynamic Mul-
 165 tiQC report (Ewels et al. 2020).

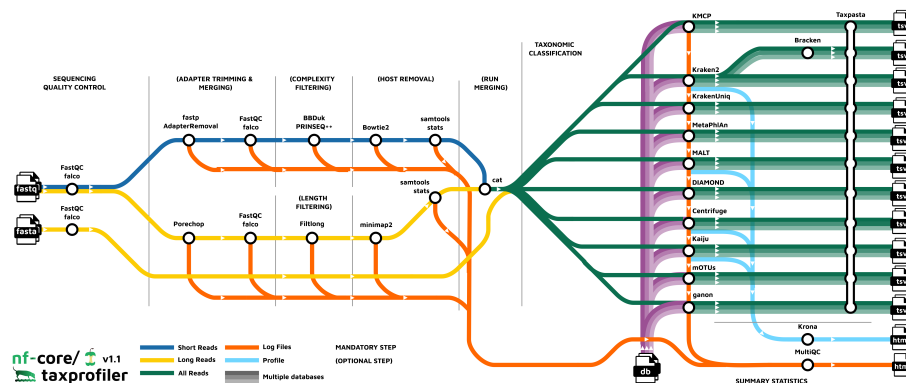


Figure 1: Visual overview of the nf-core/taxprofiler workflow. nf-core/taxprofiler can take in FASTQ (short or long reads) or FASTA files (long reads), that will optionally go through sequencing quality control (e.g. with FastQC), read preprocessing (e.g. removal of adapters), complexity filtering, host removal, and run merging before performing taxonomic classification and/or profiling with a user-selected range of tools and databases. Output from all classifiers and profilers are standardised into a common taxon table format, and when supported visualisations of the profiles are generated.

Table 1: List of nf-core/taxprofiler supported taxonomic/classifiers profilers as of version 1.1 and their approximate method and supported input database types. Primary algorithm refers to the algorithm type used for sequencing matching. Reference type refers to the typical sequence type used in database construction of the tool. Method refers to whether the tool performs just read classification (classifier) or additionally abundance estimation (profiler). Sequencing matching type refers to which ‘molecular alphabet’ is primarily used for matching between a query (read) and a reference (genome/gene).

Tool	Primary Algorithm	Reference Type	Method	Sequence Matching Type
Kraken2	k-mer based	whole-genome	classifier	Nucleotide
Kaiju	k-mer based	whole-genome	classifier	Amino Acid

Tool	Primary Algorithm	Reference Type	Method	Sequence Matching Type
Bracken	k-mer based	whole-genome	profiler	Nucleotide
KrakenUniq	k-mer based	whole-genome	profiler	Nucleotide
ganon	k-mer based	whole-genome	profiler	Nucleotide
KMCP	k-mer based	whole-genome	profiler	Nucleotide
MALT	alignment based	whole-genome	classifier	Nucleotide/Amino Acid
DIAMOND	alignment based	whole-genome	classifier	Amino Acid
Centrifuge	alignment based	whole-genome	profiler	Nucleotide
MetaPhlAn	alignment based	marker-gene	profiler	Nucleotide
mOTUS	alignment based	marker-gene	profiler	Nucleotide

166 nf-core/taxprofiler comes with extensive documentation for general usage, short- and
 167 long- parameter help texts, and output file descriptions. To ensure maximum accessi-
 168 bility, these are available in pipeline results as markdown files ([https://github.com/nf-
 169 core/taxprofiler](https://github.com/nf-core/taxprofiler)), on the nf-core website (<https://nf-co.re/taxprofiler>) and for the pa-
 170 rameter help texts on the command line via standard --help. The output documen-
 171 tation also aims to guide users as the most suitable files for different types of down-
 172 stream analysis

173 4 Discussion

174 A range of pipelines already exists for taxonomic profiling, however, each have
 175 their own particular purpose and capabilities. We compared the functionality
 176 of nf-core/taxprofiler against four other recently published or released pipelines,
 177 selected based on their similarity of purpose to nf-core/taxprofiler. The selection
 178 criteria and a more detailed comparison between the five pipelines can be seen
 179 in the [Supplementary Information](#). Overall, while there was a general similarity
 180 across all pipelines, nf-core/taxprofiler showed the largest number of options for
 181 pipeline execution accessibility, and user choice. This is facilitated through the
 182 use of an established workflow manager (with Nextflow supporting 7 software
 183 environment/container systems), support for both CLI and GUI execution, and by the
 184 number of supported classifiers. Furthermore, it is unique in that is the only pipeline
 185 to support supplying multiple database for all of the tools in a single pipeline run.

Table 2: Comparison of functionality with four recent taxonomic pipelines with similar functionality. A more detailed textual comparison can be found in the [Supplementary Information](#). Category keys are as follows: I - Information, R - Reproducibility, A - Accessibility, P - Portability, S - Scalability, F - Functionality.

Category	Criterion	StaG-mwc	sunbeam	Unipro UGENE	tama	nf-core/taxprofiler
I	Source code URL	https://github.com/ctmrbio/stag-mwc	https://github.com/sunbeam-labs/sunbeam	https://github.com/ugeneunipro/ugene	https://github.com/kimlab/TAMA	https://github.com/nf-core/taxprofiler/
I	Evaluated version	0.7.0	4	48	githash: 3a22c8f	1.1.0
I	Last release date	2023-06-13	2023-08-08	2023-08-08	2022-03-02	2023-09-19
I	Publication year	Unpublished	2019	2019	2020	This publication
I	Publication DOI	Unpublished	10.1186/s40168-019-0658-x	10.1093/bioinformatics/bt184	10.1093/bioinformatics/bt184	This publication
R	Pipeline versioning	Yes	Yes	Yes	No	Yes
R	Software versioning	Yes	Yes	Yes	Yes	Yes
R	Nr. software environments or container engines supported	2	2	0	1	7
A	Installation documentation	Yes	Yes	Yes	Yes	Yes
A	Usage documentation	Yes	Yes	Yes	Yes	Yes
A	Output documentation	Yes	Yes	Yes	Yes	Yes
A	CLI execution interface	Yes	Yes	No	Yes	Yes
A	GUI execution interface	No	No	Yes	No	Yes

Category	Criterion	StaG-mwc	sunbeam	Unipro UGENE	tama	nf-core/taxprofiler
A/S	Integration a scheduling systems	Yes	Yes	No	No	Yes
P/A	Nr. supported operating systems	2	1	3	1	2
P	Local machine integration	Yes	Yes	Yes	Yes	Yes
P/S	HPC scheduler integration	Yes	Yes	No	No	Yes
P/S	Cloud computing integration	Unsure	Unsure	No	No	Yes
P/S	Integration with multiple scheduling systems	Partial	Partial	No	No	Yes
S	Per-process resource optimisation	Yes	Yes	Yes	No	Yes
F	Short read support	Yes	Yes	Yes	Yes	Yes
F	Long read support	No	No	Yes	No	Yes
F	Read preprocessing	Yes	Yes	Yes	Yes	Yes
F	Sequencing depth estimation	Yes	No	No	No	No
F	Complexity filtering	No	Yes	No	No	Yes
F	Host removal	Yes	Yes	Partial	No	Yes
F	Nr. supported taxonomic classifiers/profilers	7	3	3	3	11
F	Graphical run reports	Yes	No	No	No	Yes
F	Standardised profiles	No	No	No	Yes	Yes

Category	Criterion	StaG-mwc	sunbeam	Unipro UGENE	tama	nf-core/taxprofiler
F	Multiple database supported	Partial	No	No	No	Yes
F	Metagenomic assembly	No	Yes	No	No	No
F	Visualisation	No	No	No	No	Partial

Another important advantage of nf-core/taxprofiler is that it is being developed within the nf-core community (<https://nf-co.re>), that provides strong long-term support for the continued community-based development and maintenance of its pipelines. In this framework, we will continue to add additional preprocessing, metagenomic classification, and profiling tools as they become established and as requested by the metagenomics community. For example, we feel that the inclusion of steps such as sequencing saturation estimation as already being performed by a similar pipeline StaG-mwc (<https://github.com/ctmrbio/stag-mwc>) would be beneficial to the nf-core/taxprofiler workflow (possibly with dedicated tools such as Nonpareil, Rodriguez-R et al. 2018), and/or more performant complexity filtering tools such as Komplexity as offered by the sunbeam metagenomics pipeline (Clarke et al. 2019). Additional tools that could be added for short-read classification could include sourmash (Titus Brown and Irber 2016) that provides scalable sequence to sequence comparison or other marker gene reference tools such as tools such as METAXA2 (Bengtsson-Palme et al. 2015) that use shotgun sequencing reads to recover 16S sequences from metagenomic samples. Adding additional classifiers also applies to extend support to other sequencing platforms; nf-core/taxprofiler already supports Nanopore long-read data, however the use of long-read PacBio data for metagenomic data is growing in interest (Portik, Brown, and Pierce-Ward 2022). We are therefore considering adding dedicated preprocessing steps for this type of sequencing data.

A remaining major challenge for metagenomics researchers (and not supported in the same workflow by any of the compared pipelines above) is the construction of databases for each profiling tool. Given there still are no curated, high-quality ‘gold standard’ databases in metagenomics, and while nf-core/taxprofiler allows the profiling against multiple databases and settings in parallel, currently the pipeline still requires users to construct these manually and to supply to the pipeline. While we feel this is currently a reasonable investment as such databases are typically repeatedly re-used, we are exploring the possibility to add an additional complementary workflow in the pipeline to allow automated database construction of all classification tools, given a set of FASTA reference files.

Finally, once an overall taxonomic profile is generated, researchers often wish to validate hits through more sensitive and accurate methods such as with read-mapping alignment. While read alignment is supported by other pipelines such as StaG-mwc,

220 this happens in-parallel to the taxonomic profiling and requires prior expectation of
221 which reference genomes to map against. Instead, nf-core/taxprofiler could be eas-
222 ily extended to have a validation step similar to the approach of the ancient DNA
223 metagenomic pipeline aMeta (Pochon et al. 2022). Utilising Nextflow’s execution par-
224 allelism, the input sequences could be aligned back to the reference genomes of only
225 those species with hits resulting from the taxonomic classification, but with dedicated
226 accurate short- or long-read aligners. In addition to the more precise classification,
227 post-classification read-alignment could also be particularly useful for researchers in
228 palaeogenomics who wish to use tools other than KrakenUniq for initial classification
229 (as in aMeta), where alignment information can be used to authenticate ancient DNA
230 within their samples, but also in clinical metagenomics to identify potential pathogens
231 at much finer resolution (e.g. down to strain level).

232 Another motivation for developing nf-core/taxprofiler, despite the large number of ex-
233 isting metagenomics pipelines, is that by establishing a taxonomic profiling pipeline
234 within the nf-core ecosystem, it is possible to begin building both standalone but
235 also an integrated suite of powerful interconnected pipelines for the major stages
236 of metagenomic workflows. Existing microbial- and metagenomics- related pipelines
237 within the nf-core initiative include nf-core/ampliseq (Straub et al. 2020), nf-core/mag
238 (Krakau et al. 2022), and nf-core/funcscan (<https://nf-co.re/funcscan>). We expect over
239 time the ability to link inputs and outputs of each workflow to develop comprehensive
240 metagenomic analyses, while still maintaining powerful standalone pipelines, provid-
241 ing maximal user choice but with familiar interfaces.

242 5 Conclusion

243 nf-core/taxprofiler is an accessible, efficient, and scalable pipeline for metagenomic
244 taxonomic classification and profiling that can be executed on anywhere from laptops
245 to the cloud. To our knowledge, the pipeline offers the largest number of taxonomic
246 profilers across similar pipelines, providing flexibility for users not just on choice of
247 profiling tool but also with databases and database settings within a single run. With
248 the development within the open and welcoming nf-core community and with best-
249 practise development infrastructure, we look forward to further contributions and in-
250 volvement of the wider metagenomics community, and also we hope that through de-
251 tailed documentation and a range of execution options, nf-core/taxprofiler will make
252 reproducible and high-throughput metagenomics more accessible for a wide range of
253 disciplines.

254 6 Code Availability

255 nf-core/taxprofiler source code is available on GitHub at [https://github.com/nf-core/](https://github.com/nf-core/taxprofiler)
256 [taxprofiler](https://github.com/nf-core/taxprofiler), and each release is archived on Zenodo (latest version DOI: [10.5281/zen-](https://doi.org/10.5281/zenodo.7728364)
257 [odo.7728364](https://doi.org/10.5281/zenodo.7728364))

258 The version of the pipeline described in this paper is version 1.1.0 (release specific

259 Zenodo archive DOI: [10.5281/zenodo.8358147](https://doi.org/10.5281/zenodo.8358147))

260 **7 Acknowledgments**

261 We thank Prof. Christina Warinner and the Microbiome Sciences group MPI-EVA for
262 original discussions that lead to the pipeline. We are also grateful for the nf-core
263 community for the original and ongoing support in the development in the pipeline, in
264 particular for the contributions by Lauri Mesilaakso, Jianhong Ou, and Rafał Stepień.

265 **8 Funding**

266 S.S. and L.A-L. were supported by Rapid establishment of comprehensive laboratory
267 pandemic preparedness – RAPID-SEQ. This material is based upon work supported by
268 the U.S. Department of Agriculture, Agricultural Research Service, under agreement
269 No. 58-3022-0-001 (T.A.C II). M.B. and J.A.F.Y were supported by the Max Planck So-
270 ciety. J.A.F.Y was supported by the Werner Siemens-Stiftung (“Paleobiotechnology”,
271 Awarded to Prof. Pierre Stallforth and Prof. Christina Warinner).

272 **9 Supplementary Information**

273 **9.1 Implementation**

274 **9.1.1 Input and Execution**

275 The pipeline can be executed via typical Nextflow commands (Code Block 1), or us-
276 ing the standard nf-core ‘launch’ GUI (Figure 2), making the pipeline accessible for
277 both computationally experienced as well as less experienced researchers. In addi-
278 tion to the general usage and parameter documentation of the pipeline ([https://nf-
279 co.re/taxprofiler](https://nf-co.re/taxprofiler)). The GUI offers immediate assistance and guidance to users on what
280 each parameter does, both in short- and long-form, with long-form parameter descrip-
281 tions additionally describing which tool-specific parameters are being modified for
282 each pipeline parameter (<https://nf-co.re/taxprofiler/launch>). The GUI also includes
283 controlled user input by providing strict drop-down lists and input validation prior
284 execution of the pipeline (Figure 2) to reduce the risk of typos and other mistakes,
285 which is in contrast to the command-line interface that only includes validation at
286 pipeline run-time.

287 An example nf-core command line execution of the pipeline can be seen in Code
288 Block 1, where two input files are supplied: one file specifying paths of FASTQ files
289 of metagenomic samples and necessary metadata for preprocessing (such as sample
290 ID and sequencing platform), and the second file specifying paths to the user-defined
291 databases with per-database classification parameters. Various parameters are avail-
292 able to select different preprocessing steps, and provide additional configuration such
293 as tool selection and value options. Note that even if a user supplies a given database
294 in the database input sheet, the corresponding profiling tool must still be activated

Preprocessing short-read QC options

Launch

--shortread_qc_minlength

15

?

Specify the minimum length of reads to be retained

Specifying a minimum read length filtering can speed up profiling by reducing the number of short unspecific reads that need to be match/aligned to the database.

Modifies tool parameter(s):

- removed from reads --length_required
- AdapterRemoval: --min length

--perform_shortread_complexityfilter

☐ True
☒ False

?

Turns on nucleotide sequence complexity filtering

--shortread_complexityfilter_tool

bbduk

?

Specify which tool to use for complexity filtering

[Select an option]

bbduk

prinseqplusplus

fastp

--shortread_complexityfilter_entropy

?

Specify the minimum sequence entropy level for complexity filtering

--shortread_complexityfilter_bbduk_windowsize

50

?

On this page

Nextflow command-line flags

> Input/output options

Preprocessing general QC options

Preprocessing short-read QC options

Preprocessing long-read QC options

Preprocessing host removal options

Preprocessing run merging options

Profiling options

Postprocessing and visualisation options

Show hidden params

Figure 2: Screenshot of the nf-core pipeline launch graphical user interface with nf-core/taxprofiler options displayed. The web browser-based interface provides guidance for how to configure each pipeline parameter by providing both short and long help descriptions to help guide users in which contexts to configure each parameter. Additional elements such as radio buttons, drop down menus, and background regular expressions check for validity of input. When pressing launch, a prepared configuration file and command is provided that can be copied and pasted by the user into the terminal

Listing 1 Example nf-core/taxprofiler command for running short-read quality control, removal of host DNA and executing the k-mer based Kraken2 and marker gene alignment MetaPhlAn tools.

```
$ nextflow run nf-core/taxprofiler \  
-r 1.1.0 \  
-profile singularity,<institute> \  
--input <samplesheet.csv> \  
--databases <database.csv> \  
--perform_shortread_qc \  
--shortread_qc_minlength 20 \  
--preprocessing_qc_tool falco \  
  
--run_host_removal --hostremoval_reference 'host_genome.fasta' \  
--run_kraken2 --kraken2_save_reads \  
--run_metaphlan \  
--run_krona \  
--run_profile_standardisation
```

295 with the corresponding pipeline parameter (e.g. --run_kraken2). Per-classifier flags
296 are also available for the optional saving of additional non-profile output files. Alter-
297 natively to command line flags, parameters can be specified via pre-configured YAML
298 format files, with which (provided no hardcoded paths are included) can be re-used
299 across pipeline runs.

300 All nf-core pipelines are strictly versioned (specified with the Nextflow -r flag), and to
301 ensure reproducibility, each version of the pipeline has a fixed set of software used for
302 each step of the pipeline. The fixed set of software are controlled through the use of
303 the conda package manager or containers (Docker, or Apptainer - previously known
304 as Singularity, etc) from the stable Bioconda (Grüning et al. 2018) or BioContainers
305 (Veiga Leprevost et al. 2017) repositories. This, coupled with the intrinsic Nextflow
306 ability to execute on most infrastructure whether that is a local laptop (resource re-
307 quirements permitting), traditional HPC, as well across common cloud providers also
308 makes nf-core/taxprofiler a very portable pipeline that can be used in many contexts.

309 9.1.2 Preprocessing

310 Preprocessing steps in nf-core/taxprofiler are aimed at removing laboratory and se-
311 quencing artefacts that may influence taxonomic profiling, either for computing re-
312 source consumption or and/or false-positive or false-negative classification reasons.
313 First sequencing quality control with FastQC (Andrews 2010) or Falco (Sena Brandine
314 and Smith 2021) is carried out. Falco was included for reduced memory requirements,
315 in particular for long read sequencing data. Artificial library adapter sequences added
316 during sequencing reduce sequencing matching accuracy by reducing sequence speci-

317 ficity, and in some cases, may result in false-positive hits due to adapter sequence con-
 318 tamination in reference genomes (Schäffer et al. 2018; F. P. Breitwieser, Baker, and
 319 Salzberg 2018)¹. Additionally, paired-end merging may provide longer sequences
 320 that will allow for more specific classification when paired-end alignment is not sup-
 321 ported by a given classifier. For these tasks nf-core/taxprofiler can apply either fastp
 322 (Chen et al. 2018) or AdapterRemoval2 (Schubert, Lindgreen, and Orlando 2016) for
 323 short reads, and currently Porechop (Wick et al. 2017) for Oxford Nanopore long-read
 324 data. For both short and long reads, FastQC or Falco is run again to allow assessment
 325 on the performance of the adapter removal and/or pair-merging step.

326 Low complexity sequences, e.g. sequences containing long stretches of mono- or
 327 di-nucleotide repeats provide little specific genetic information that contribute to
 328 taxonomic identification, as they can align to many different reference genomes
 329 (Schmieder and Edwards 2011; Clarke et al. 2019). Including such reads during
 330 taxonomic profiling can increase run-time and memory usage for little gain, as
 331 during lowest-common-ancestor (LCA) classification steps they will be assigned to
 332 high-level taxonomic ranks (e.g. Kingdom). nf-core/taxprofiler performs removal of
 333 these reads through complexity filtering algorithms as provided by fastp, BBDuk
 334 (Bushnell 2022), or PRINSEQ++ (Cantu, Sadural, and Edwards 2019). Long read
 335 sequences often do not have such reads, as lengths are sufficient enough to capture
 336 greater sequence diversity - but it is sometimes desirable to only classify reads longer
 337 than a certain length - as these provide more precise taxonomic information (Dilthey
 338 et al. 2019; Portik, Brown, and Pierce-Ward 2022). Therefore, nf-core/taxprofiler can
 339 remove reads shorter than a user-defined length using Filtlong.

340 Removing host DNA is another common preprocessing step in metagenomic studies.
 341 This can help speed up run-time, particularly in microbiome studies, where detection
 342 of microbes are of interest. Furthermore, host-contamination of reference genomes in
 343 public databases is common (Longo, O’Neill, and O’Neill 2011; Kryukov and Imanishi
 344 2016; Florian P. Breitwieser et al. 2019). Therefore, the removal of such sequences can
 345 help decrease the risk of false positive taxonomic assignment. To remove multiple
 346 hosts or other sequences, all reference genomes can be combined into a single FASTA
 347 reference file. Short read host removal can be carried out with Bowtie2 (Langmead
 348 and Salzberg 2012; Langmead et al. 2019) and minimap2 (Li 2018) for long reads, both
 349 in combination with SAMtools (Li et al. 2009; Danecek et al. 2021), where reads are
 350 aligned against the reference genome and the off-target (unaligned) reads are then
 351 converted back to FASTQ format for classification.

352 Finally, nf-core/taxprofiler can optionally perform ‘run merging’ where multiple
 353 FASTQ files from the same sample but have been sequenced over multiple lanes are
 354 concatenated together to generate one profile per sample or library. The final set of

¹For an ‘infamous’ case of adapter sequences in a published eukaryotic genome, see the following blog posts

Graham Etherington: <https://web.archive.org/web/20201219022000/http://grahametherington.blogspot.com/2014/09/why-you-should-qc-your-reads-and-your.html?m=1> why-you-should-qc-your-reads-and-your.html
 Sixing Huang: <https://web.archive.org/web/20220904205331/https://dgg32.medium.com/carp-in-the-soil-1168818d2191>

(Accessed 2023-08-25)

reads used for profiling can be optionally saved for downstream re-use. Throughout all steps, relevant statistics and log files are generated and used both for the final pipeline run report as well as saved into the results directory of the pipeline run for further inspection where necessary.

9.1.3 Profiling

There are many types of metagenomic profiling techniques, from profiling against whole-genome references with alignment or k-mer based approaches, to methods involving alignment to species-specific marker-gene families (Quince et al. 2017; Ye et al. 2019). `nf-core/taxprofiler` aims to support and include all established classification or profiling tools as requested by the community.

The choice of tools used in a pipeline run is up to the user, with a tool being executed when both the corresponding database and `--run_<tool>` flag is provided. Specific classification settings for each tool and database are specified in the database CSV input sheet. Some tools also have pipeline level command-line flags for controlling certain aspects of output files.

The following classifiers and profilers are supported in version 1.1.0 of `nf-core/taxprofiler`: Kraken2 (Wood, Lu, and Langmead 2019), Bracken (Lu et al. 2017), KrakenUniq (F. P. Breitwieser, Baker, and Salzberg 2018), Centrifuge (Kim et al. 2016), MALT (Vågene et al. 2018), DIAMOND (Buchfink, Reuter, and Drost 2021), Kaiju (Menzel, Ng, and Krogh 2016), MetaPhlAn (Blanco-Míguez et al. 2023), mOTUs (Ruscheweyh et al. 2022), ganon (Piro et al. 2020), KMCP (Shen et al. 2023).

By default, `nf-core/taxprofiler` produces the default per-sample taxonomic classification profile output from a tool or a tool's report generation tool. The output is normally in the form of counts per reference sequencing, with additional statistics about the hits of a particular organism (estimated abundance, taxonomic level etc.). Users can also optionally request output of per-read classification output and output such as classified and unclassified reads in FASTQ format, where supported.

The pipeline provides high efficiency, particularly during the metagenomic classification stage, through the inherent parallelisation provided by Nextflow. While metagenomic classification is comparatively computationally intensive (in terms of memory and execution time; due to a combination of sequencing depth and number of reference genomes), Nextflow automatically optimises the execution order of all the steps in pipeline, maximising the number parallel running of multiple profilers and/or databases at any given time point, as far as the available computational resources allow. For local machines such as laptops or desktops, Nextflow will automatically detect all available computational resources, but this is customisable using Nextflow configuration files. For HPC and cloud infrastructure, users typically have to define the computational infrastructural environment the pipeline is being executed on (CPU or memory limitations, queues, instance types, etc.). To facilitate the pipeline computational configuration, `nf-core/taxprofiler` supports use of more than 90 pre-defined centralised generic and pipeline-specific institutional Nextflow configurations as provided by `nf-core/configs` (<https://nf-co.re/configs>). However, of course users are still

397 welcome to supply their own custom configuration files as with any typical Nextflow
398 run, further refining computational limitations or execution specifications.

399 **9.1.4 Post-profiling**

400 In metagenomic studies, it is common practise to compare the profiles among many
401 samples, and the results of multiple profiles are normally stored in ‘taxon tables’, i.e.,
402 counts per reference taxon (rows), for each sample (columns). When available, nf-
403 core/taxprofiler supports the option to produce the ‘native’ taxon table of each classi-
404 fication tool when multiple samples are run.

405 One of the challenges that researchers face when comparing multiple taxonomic clas-
406 sifiers or profilers is the heterogenous output formats that are produced, that often
407 require custom parsing and merging scripts for each tool to standardise. To facilitate
408 more user-friendly cross-comparisons between tools, nf-core/taxprofiler utilises the
409 TAXPASTA tool (Beber et al. 2023) to generate standardised profiles and generate
410 multi-sample tables.

411 Summary statistics for the entire pipeline are visualised and displayed in a customis-
412 able MultiQC report (Ewels et al. 2020). When supported, quality control of data and
413 pipeline runs are shown for manual verification. Krona plots (Ondov, Bergman, and
414 Phillippy 2011) can also optionally be generated for supported tools to help provide
415 further visualisation of taxonomic profiles.

416 **9.1.5 Output**

417 To summarise, the main default output from nf-core/taxprofiler are both classifier
418 ‘native’ and standardised single- and multi-sample taxonomic profiles with counts
419 per-taxon and an interactive MultiQC run report with all run statistics, in addition to
420 the raw log files themselves where available.

421 The MultiQC run report displays statistics and summary visualisations for all steps of
422 the pipeline where possible, lists of versions for all tools of each step of the pipeline.
423 It also provides a dynamically-constructed text for the recommended ‘methods’ for
424 reporting how the pipeline was executed (including relevant citations) that users can
425 use in their own publications.

426 Optional outputs can include other types of profiles (e.g. per read classification) and
427 in other formats as produced by the tools themselves, as well as raw reads from pre-
428 processing steps and output visualisations from Krona. Nextflow resource usage and
429 trace reports are also by default produced for users to check pipeline performance.

430 **9.2 Comparison with other solutions**

431 nf-core/taxprofiler has been specifically developed for the analysis of whole-genome,
432 *metagenomic* sequencing data. While other types of taxonomic profiling data such
433 as 16S amplicon sequencing are well established fields with a range of popular high-
434 quality and best-practise tools pipelines (e.g. Blanco-Míguez et al. 2023; Schloss et

al. 2009) and databases (DeSantis et al. 2006; Yilmaz et al. 2014), ‘gold standard’ tools and databases for metagenomics remain much less established. Thus, the need for highly-multiplexed classification is more desirable for the newer metagenomics methods.

We searched Google Scholar for open-source pipelines published or released in the last 5 years (at the time of writing, since 2018) that were designed primarily for metagenomic classification screening, that supported at least 2 classifiers, had at least one preprocessing step and were not specifically targeted at read classification of specific domains of taxa (e.g. viruses or bacteriophages only). We also included an additional open-source but unpublished pipeline at the recommendations of the authors of the pipeline due to the functional overlap to nf-core/taxprofiler. We then evaluated the pipelines based on their publications and documentation for typical metagenomic profiling workflow steps. We used a range of criteria related to expectations of modern bioinformatic workflows that can be summarised in the following four categories: reproducibility, accessibility, scalability, and portability (Wratten, Wilm, and Göke 2021). After searching, we selected the following pipelines for comparison with nf-core/taxprofiler that matched the specific criteria described above: sunbeam (v4, Clarke et al. 2019), Unipro UGENE (v48, Rose et al. 2019), TAMA (github: 3a22c8f, Sim et al. 2020), and StaG-mwc (0.7.0, Boulund et al. 2023).

In terms of accessibility, all pipelines have documentation describing the installation steps, usage instructions, and output files. However, there are varying levels of detail and comprehensiveness. In particular, StaG-mwc and nf-core/taxprofiler have the most detailed descriptions of all possible output files for every supported module, whereas Unipro UGENE and sunbeam have very minimal to possibly unfinished output documentation. For execution options, most of the pipelines provide CLI execution, except for Unipro UGENE which offers only GUI-based pipeline set-up (despite a command-line execution of the GUI generated configuration). In particular, nf-core/taxprofiler is the only pipeline providing both CLI and GUI interfaces for pipeline run execution.

Criteria covering portability also overlap with accessibility, as it implies options for and ease of different users running on different types of computing infrastructure, whether that is on their own laptop, on an HPC cluster, or in the cloud. Unipro UGENE is the only pipeline that explicitly states support for execution on all three major operating systems (Linux, OSX, Windows), whereas StaG-mwc and nf-core/taxprofiler can be run on unix operating systems (albeit possibly on Windows via Windows Subsystem for Linux (WSL)), and sunbeam and TAMA are only being supported on Linux.

While all pipelines support ‘local’ machine execution (e.g. personal laptops or desktops), a large portion of academic users execute computationally intensive bioinformatic tasks on HPC clusters. In these contexts, pipeline task submissions are normally managed by job schedulers, thus integration with schedulers is an important criterion for running large multi-step and parallelised pipelines. The three pipelines leveraging workflow managers (Snakemake and Nextflow) support integration with schedulers (StaG-mwc, sunbeam, and nf-core/taxprofiler) with nf-core/taxprofiler supporting the most by far ([>10 scheduling systems](#)) as natively offered by Nextflow. This allows

the greatest possible choice for users in terms of which HPC infrastructure they can execute their pipeline on. As an extension of this, only nf-core/taxprofiler has explicit support for cloud computing (e.g. AWS, GCP, or Microsoft Azure) as provided by Nextflow, again maximising user choice and portability when it comes to running the pipeline.

In terms of scalability, the aforementioned integration with schedulers and cloud computing support implicitly maximises efficiency and parallelisation of pipeline runs, providing good scalability for varying numbers of input files and steps in the pipeline. Again, the three workflow manager based pipelines provide scalability, whereas there is no mention neither Unipro UGENE nor TAMA in reference to parallel task execution. Furthermore, all pipelines except TAMA, allowed per-process customisation of computational resources, something critical for maximising efficient scalability to ensure only the necessary resources for a given step of a pipeline are requested.

In terms of reproducibility, all five pipelines are good at ensuring reproducibility in terms of pipeline and software versioning (allowing re-execution of pipeline runs using the same software), with only TAMA not having stable versioned releases. However, installing software manually across different infrastructures can result in variability in the execution of each software² (Di Tommaso et al. 2017). The current most popular solution to the problem of inconsistent software environments is to use container engines such as Docker or Apptainer to run container images which are isolated, deterministic computing environments which can be executed by any system providing a container runtime. Only Unipro UGENE does not document the use of a container system, with nf-core/taxprofiler offering the biggest choice for users, again, courtesy of Nextflow with 6 different engine systems at the time of writing.

Finally, we compared metagenomics related functionality between the pipelines. All pipelines support short-read FASTQ input, but only nf-core/taxprofiler explicitly reports long-read support, while the documentation in Unipro UGENE states that assembled contigs are possible input to some of the profilers. All pipelines support read preprocessing (adapter clipping, and merging). In terms of tools used for preprocessing, Trimmomatic (Bolger, Lohse, and Usadel 2014) is popular across the other pipelines but is not supported in nf-core/taxprofiler. Only sunbeam and nf-core/taxprofiler support complexity filtering to remove low sequence diversity reads. In fact within sunbeam, the authors developed their own dedicated, performant complexity filtering tool Komplexity (Clarke et al. 2019). Most pipelines support some form of host removal (only TAMA did not support this), and it is likely possible with Unipro UGENE (although not directly described). In all cases, host removal consists of mapping processed reads with an aligner and using the off-target reads for downstream profiling (as implemented in nf-core/taxprofiler), however StaG-mwc has an additional separate metagenomic host removal step with Kraken2. nf-core/taxprofiler supports by far the largest number of taxonomic classifiers and profilers at 11 as of v1.1.0 - providing the greatest choice to users - with StaG-mwc offering 7, and the remaining

²As demonstrated in this blogpost from Paweł Przytuła: <https://web.archive.org/web/20230320223436/https://appsilon.com/reproducible-research-when-your-results-cant-be-reproduced/> (Accessed 2023-08-25)

520 pipelines only 3. Only nf-core/taxprofiler and partly StaG-mwc explicitly support run-
 521 ning each profiler with multiple databases. nf-core/taxprofiler is the only pipeline that
 522 supports running an arbitrary number of different metagenomic profiler databases
 523 each with their own settings. This makes it a useful for tool parameter compari-
 524 son, testing different databases, or reducing the size of each database (e.g. per do-
 525 main) to make it more flexibility for running on smaller computational infrastructure.
 526 StaG-mwc allows multiple references for their short-read alignment steps rather than
 527 the metagenomic profilers. For output, nf-core/taxprofiler, StaG-mwc, and sunbeam
 528 (via an extension) support a singular run report for summarising all preprocessing
 529 step. Only nf-core/taxprofiler and TAMA produce standardised output for all taxo-
 530 nomic profilers, the former with the dedicated standalone tool TAXPASTA (Beber et
 531 al. 2023). However Unipro UGENE additionally offers a ‘consensus’ profile using
 532 WEVOTE (Metwally et al. 2016).

533 To summarise, many of the pipelines reviewed here offer similar functionality, with
 534 particularly StaG-mwc having a strong overlap with nf-core/taxprofiler. Thus, users
 535 in most cases will be able to select the pipeline depending on which framework they
 536 feel most comfortable with. However the advantages of nf-core/taxprofiler mainly
 537 come from the offering of the greatest choice of tools, as well the particular benefits
 538 provided by Nextflow. It provides the greatest number of computational infrastruc-
 539 ture types the pipeline can be executed on, and container systems can be used to
 540 ensure reproducibility, as well the support of the nf-core community due to the cen-
 541 tralised pool of ‘plug-and-play’ modules to make it easier to update the pipeline over
 542 time to add new tools classifiers.

543 The functionality offered by other pipelines not currently supported by nf-
 544 core/taxprofiler include sequencing saturation estimation (StaG-mwc), taxonomy-
 545 free composition comparison (StaG-mwc), functional profiling (StaG-mwc), *de novo*
 546 assembly (sunbeam), and reference mapping (StaG-mwc, sunbeam). We do not
 547 plan to support *de novo* assembly or functional profiling in nf-core/taxprofiler as
 548 we feel these are already better served by other existing dedicated pipelines within
 549 the nf-core ecosystem: nf-core/mag for *de novo* assembly, (Krakau et al. 2022)
 550 and nf-core/funcscan for functional profiling (<https://nf-co.re/funcscan>), as well as
 551 elsewhere e.g. MetaWRAP (Uritskiy, DiRuggiero, and Taylor 2018).

552 References

- 553 Andrews, Simon. 2010. “FastQC: A Quality Control Tool for High Throughput Se-
 554 quence Data.” <http://www.bioinformatics.babraham.ac.uk/projects/fastqc/>.
- 555 Beber, Moritz E, Maxime Borry, Sofia Stamouli, and James A Fellows Yates. 2023.
 556 “TAXPASTA: TAXonomic Profile Aggregation and STAndardisation.” *Journal of*
 557 *Open Source Software* 8 (87): 5627. <https://doi.org/10.21105/joss.05627>.
- 558 Bengtsson-Palme, Johan, Martin Hartmann, Karl Martin Eriksson, Chandan Pal, Kaisa
 559 Thorell, Dan Göran Joakim Larsson, and Rolf Henrik Nilsson. 2015. “METAXA2:
 560 Improved Identification and Taxonomic Classification of Small and Large Subunit
 561 rRNA in Metagenomic Data.” *Molecular Ecology Resources* 15 (6): 1403–14. <https://doi.org/10.1111/1365-3113.12403>.

562 [//doi.org/10.1111/1755-0998.12399](https://doi.org/10.1111/1755-0998.12399).

563 Blanco-Míguez, Aitor, Francesco Beghini, Fabio Cumbo, Lauren J McIver,
564 Kelsey N Thompson, Moreno Zolfo, Paolo Manghi, et al. 2023. "Extend-
565 ing and Improving Metagenomic Taxonomic Profiling with Uncharacter-
566 ized Species Using MetaPhlAn 4." *Nature Biotechnology*, February, 1–12.
567 <https://doi.org/10.1038/s41587-023-01688-w>.

568 Bolger, Anthony M, Marc Lohse, and Bjoern Usadel. 2014. "Trimmomatic: A Flexible
569 Trimmer for Illumina Sequence Data." *Bioinformatics (Oxford, England)* 30 (15):
570 2114–20. <https://doi.org/10.1093/bioinformatics/btu170>.

571 Boulund, Fredrik, Aron Arzoomand, Justine Debelius, chrsb, and Lisa Olsson. 2023.
572 "Ctmbio/Stag-Mwc: StaG v0.7.0." Zenodo. [https://doi.org/10.5281/ZENODO.](https://doi.org/10.5281/ZENODO.8032462)
573 [8032462](https://doi.org/10.5281/ZENODO.8032462).

574 Breitwieser, F P, D N Baker, and S L Salzberg. 2018. "KrakenUniq: Confident and Fast
575 Metagenomics Classification Using Unique k-Mer Counts." *Genome Biology* 19 (1):
576 198. <https://doi.org/10.1186/s13059-018-1568-0>.

577 Breitwieser, Florian P, Jennifer Lu, and Steven L Salzberg. 2019. "A Review of Meth-
578 ods and Databases for Metagenomic Classification and Assembly." *Briefings in*
579 *Bioinformatics* 20 (4): 1125–36. <https://doi.org/10.1093/bib/bbx120>.

580 Breitwieser, Florian P, Mihaela Pertea, Aleksey Zimin, and Steven L Salzberg. 2019.
581 "Human Contamination in Bacterial Genomes Has Created Thousands of Spurious
582 Proteins." *Genome Research* 29 (May): 954–60. [https://doi.org/10.1101/gr.245373.](https://doi.org/10.1101/gr.245373.118)
583 [118](https://doi.org/10.1101/gr.245373.118).

584 Buchfink, Benjamin, Klaus Reuter, and Hajk-Georg Drost. 2021. "Sensitive Protein
585 Alignments at Tree-of-Life Scale Using DIAMOND." *Nature Methods* 18 (4): 366–
586 68. <https://doi.org/10.1038/s41592-021-01101-x>.

587 Bushnell, Brian. 2022. "BBMap." <https://sourceforge.net/projects/bbmap/>.

588 Cantu, Vito Adrian, Jeffrey Sadural, and Robert Edwards. 2019. "PRINSEQ++, a Multi-
589 Threaded Tool for Fast and Efficient Quality Control and Preprocessing of Se-
590 quencing Datasets." e27553v1. PeerJ Preprints; PeerJ Inc. [https://doi.org/10.7287/](https://doi.org/10.7287/peerj.preprints.27553v1)
591 [peerj.preprints.27553v1](https://doi.org/10.7287/peerj.preprints.27553v1).

592 Chen, Shifu, Yanqing Zhou, Yaru Chen, and Jia Gu. 2018. "Fastp: An Ultra-Fast All-
593 in-One FASTQ Preprocessor." *Bioinformatics* 34 (17): i884–90. [https://doi.org/10.](https://doi.org/10.1093/bioinformatics/bty560)
594 [1093/bioinformatics/bty560](https://doi.org/10.1093/bioinformatics/bty560).

595 Chiu, Charles Y, and Steven A Miller. 2019. "Clinical Metagenomics." *Nature Reviews.*
596 *Genetics* 20 (6): 341–55. <https://doi.org/10.1038/s41576-019-0113-7>.

597 Clarke, Erik L, Louis J Taylor, Chunyu Zhao, Andrew Connell, Jung-Jin Lee, Bryton
598 Fett, Frederic D Bushman, and Kyle Bittinger. 2019. "Sunbeam: An Extensible
599 Pipeline for Analyzing Metagenomic Sequencing Experiments." *Microbiome* 7 (1):
600 46. <https://doi.org/10.1186/s40168-019-0658-x>.

601 Danecek, Petr, James K Bonfield, Jennifer Liddle, John Marshall, Valeriu Ohan, Mar-
602 tin O Pollard, Andrew Whitwham, et al. 2021. "Twelve Years of SAMtools and
603 BCftools." *GigaScience* 10 (2). <https://doi.org/10.1093/gigascience/giab008>.

604 DeSantis, T Z, P Hugenholtz, N Larsen, M Rojas, E L Brodie, K Keller, T Huber, D
605 Dalevi, P Hu, and G L Andersen. 2006. "Greengenes, a Chimera-Checked 16S
606 rRNA Gene Database and Workbench Compatible with ARB." *Applied and Envi-*
607 *ronmental Microbiology* 72 (7): 5069–72. <https://doi.org/10.1128/AEM.03006-05>.

608 Di Tommaso, Paolo, Maria Chatzou, Evan W Floden, Pablo Prieto Barja,
609 Emilio Palumbo, and Cedric Notredame. 2017. "Nextflow Enables Repro-
610 ducible Computational Workflows." *Nature Biotechnology* 35 (4): 316–19.
611 <https://doi.org/10.1038/nbt.3820>.

612 Diltthey, Alexander T, Chirag Jain, Sergey Koren, and Adam M Phillippy. 2019. "Strain-
613 Level Metagenomic Assignment and Compositional Estimation for Long Reads
614 with MetaMaps." *Nature Communications* 10 (1): 3066. <https://doi.org/10.1038/s41467-019-10934-2>.

615 Eloe-Fadrosch, Emiley A, Natalia N Ivanova, Tanja Woyke, and Nikos C Kyrpides. 2016.
616 "Metagenomics Uncovers Gaps in Amplicon-Based Detection of Microbial Diver-
617 sity." *Nature Microbiology* 1 (4): 15032. <https://doi.org/10.1038/nmicrobiol.2015.32>.

618 Ewels, Philip A, Alexander Peltzer, Sven Fillinger, Harshil Patel, Johannes
619 Alneberg, Andreas Wilm, Maxime Ulysse Garcia, Paolo Di Tommaso,
620 and Sven Nahnsen. 2020. "The Nf-Core Framework for Community-
621 Curated Bioinformatics Pipelines." *Nature Biotechnology* 38 (3): 276–78.
622 <https://doi.org/10.1038/s41587-020-0439-x>.

623 Govender, Kumeren N, and David W Eyre. 2022. "Benchmarking Taxonomic Classi-
624 fiers with Illumina and Nanopore Sequence Data for Clinical Metagenomic Diag-
625 nostic Applications." *Microbial Genomics* 8 (10): 000886. <https://doi.org/10.1099/mgen.0.000886>.

626
627
628 Grüning, Björn, Ryan Dale, Andreas Sjödin, Brad A Chapman, Jillian Rowe,
629 Christopher H Tomkins-Tinch, Renan Valieris, Johannes Köster, and Bio-
630 conda Team. 2018. "Bioconda: Sustainable and Comprehensive Soft-
631 ware Distribution for the Life Sciences." *Nature Methods* 15 (7): 475–76.
632 <https://doi.org/10.1038/s41592-018-0046-7>.

633 Hillmann, Benjamin, Gabriel A Al-Ghalith, Robin R Shields-Cutler, Qiyun Zhu, Daryl
634 M Gohl, Kenneth B Beckman, Rob Knight, and Dan Knights. 2018. "Evaluating the
635 Information Content of Shallow Shotgun Metagenomics." *mSystems* 3 (6). <https://doi.org/10.1128/mSystems.00069-18>.

636
637 Kim, Daehwan, Li Song, Florian P Breitwieser, and Steven L Salzberg. 2016. "Cen-
638 trifuge: Rapid and Sensitive Classification of Metagenomic Sequences." *Genome*
639 *Research* 26 (12): 1721–29. <https://doi.org/10.1101/gr.210641.116>.

640 Krakau, Sabrina, Daniel Straub, Hadrien Gourel, Gisela Gabernet, and Sven Nahnsen.
641 2022. "Nf-Core/Mag: A Best-Practice Pipeline for Metagenome Hybrid Assembly
642 and Binning." *NAR Genomics and Bioinformatics* 4 (1). <https://doi.org/10.1093/nargab/lqac007>.

643
644 Kryukov, Kirill, and Tadashi Imanishi. 2016. "Human Contamination in Public
645 Genome Assemblies." *PloS One* 11 (9): e0162424. <https://doi.org/10.1371/journal.pone.0162424>.

646
647 Langmead, Ben, and Steven L Salzberg. 2012. "Fast Gapped-Read Alignment with
648 Bowtie 2." *Nature Methods* 9 (4): 357–59. <https://doi.org/10.1038/nmeth.1923>.

649 Langmead, Ben, Christopher Wilks, Valentin Antonescu, and Rone Charles. 2019.
650 "Scaling Read Aligners to Hundreds of Threads on General-Purpose Processors."
651 *Bioinformatics* 35 (3): 421–32. <https://doi.org/10.1093/bioinformatics/bty648>.

652 Li, Heng. 2018. "Minimap2: Pairwise Alignment for Nucleotide Sequences." *Bioinform-*
653 *atics* 34 (18): 3094–3100. <https://doi.org/10.1093/bioinformatics/bty191>.

Li, Heng, Bob Handsaker, Alec Wysoker, Tim Fennell, Jue Ruan, Nils Homer, Gabor Marth, Goncalo Abecasis, Richard Durbin, and 1000 Genome Project Data Processing Subgroup. 2009. "The Sequence Alignment/Map Format and SAMtools." *Bioinformatics* 25 (16): 2078–79. <https://doi.org/10.1093/bioinformatics/btp352>.

Longo, Mark S, Michael J O'Neill, and Rachel J O'Neill. 2011. "Abundant Human DNA Contamination Identified in Non-Primate Genome Databases." *PloS One* 6 (2): e16410. <https://doi.org/10.1371/journal.pone.0016410>.

Lu, Jennifer, Florian P Breitwieser, Peter Thielen, and Steven L Salzberg. 2017. "Bracken: Estimating Species Abundance in Metagenomics Data." *PeerJ Computer Science* 3 (e104): e104. <https://doi.org/10.7717/peerj-cs.104>.

Lynch, Michael D J, and Josh D Neufeld. 2015. "Ecology and Exploration of the Rare Biosphere." *Nature Reviews. Microbiology* 13 (4): 217–29. <https://doi.org/10.1038/nrmicro3400>.

McIntyre, Alexa B R, Rachid Ounit, Ebrahim Afshinnekoo, Robert J Prill, Elizabeth Hénaff, Noah Alexander, Samuel S Minot, et al. 2017. "Comprehensive Benchmarking and Ensemble Approaches for Metagenomic Classifiers." *Genome Biology* 18 (1): 182. <https://doi.org/10.1186/s13059-017-1299-7>.

Menzel, Peter, Kim Lee Ng, and Anders Krogh. 2016. "Fast and Sensitive Taxonomic Classification for Metagenomics with Kaiju." *Nature Communications* 7 (April): 11257. <https://doi.org/10.1038/ncomms11257>.

Metwally, Ahmed A, Yang Dai, Patricia W Finn, and David L Perkins. 2016. "WEVOTE: Weighted VOTing Taxonomic idEntification Method of Microbial Sequences." *PloS One* 11 (9): e0163527. <https://doi.org/10.1371/journal.pone.0163527>.

Meyer, Fernando, Adrian Fritz, Zhi-Luo Deng, David Koslicki, Till Robin Lesker, Alexey Gurevich, Gary Robertson, et al. 2022. "Critical Assessment of Metagenome Interpretation: The Second Round of Challenges." *Nature Methods* 19 (4): 429–40. <https://doi.org/10.1038/s41592-022-01431-4>.

Mitchell, Alex L, Alexandre Almeida, Martin Beracochea, Miguel Boland, Josephine Burgin, Guy Cochrane, Michael R Crusoe, et al. 2019. "MGnify: The Microbiome Analysis Resource in 2020." *Nucleic Acids Research*, November. <https://doi.org/10.1093/nar/gkz1035>.

Mölder, Felix, Kim Philipp Jablonski, Brice Letcher, Michael B Hall, Christopher H Tomkins-Tinch, Vanessa Sochat, Jan Forster, et al. 2021. "Sustainable Data Analysis with Snakemake." *F1000Research* 10 (January): 33. <https://doi.org/10.12688/f1000research.29032.2>.

Morais, Diego A A, João V F Cavalcante, Shênia S Monteiro, Matheus A B Pasquali, and Rodrigo J S Dalmolin. 2022. "MEDUSA: A Pipeline for Sensitive Taxonomic Classification and Flexible Functional Annotation of Metagenomic Shotgun Sequences." *Frontiers in Genetics* 13 (March): 814437. <https://doi.org/10.3389/fgene.2022.814437>.

Nasko, Daniel J, Sergey Koren, Adam M Phillippy, and Todd J Treangen. 2018. "RefSeq Database Growth Influences the Accuracy of k-Mer-Based Lowest Common Ancestor Species Identification." *Genome Biology* 19 (1): 165. <https://doi.org/10.1186/s13059-018-1554-6>.

Nayfach, Stephen, and Katherine S Pollard. 2016. "Toward Accurate and Quantitative Comparative Metagenomics." *Cell* 166 (5): 1103–16. <https://doi.org/10.1016/j.cell.>

2016.08.007.

- Ondov, Brian D, Nicholas H Bergman, and Adam M Phillippy. 2011. "Interactive Metagenomic Visualization in a Web Browser." *BMC Bioinformatics* 12 (1): 385. <https://doi.org/10.1186/1471-2105-12-385>.
- Piro, Vitor C, Temesgen H Dadi, Enrico Seiler, Knut Reinert, and Bernhard Y Renard. 2020. "Ganon: Precise Metagenomics Classification Against Large and up-to-Date Sets of Reference Sequences." *Bioinformatics (Oxford, England)* 36 (Suppl_1): i12–20. <https://doi.org/10.1093/bioinformatics/btaa458>.
- Piro, Vitor C, Marcel Matschkowski, and Bernhard Y Renard. 2017. "MetaMeta: Integrating Metagenome Analysis Tools to Improve Taxonomic Profiling." *Microbiome* 5 (1): 101. <https://doi.org/10.1186/s40168-017-0318-y>.
- Pochon, Zoé, Nora Bergfeldt, Emrah Kırdök, Mário Vicente, Thijessen Naidoo, Tom van der Valk, N Ezgi Altınışık, et al. 2022. "aMeta: An Accurate and Memory-Efficient Ancient Metagenomic Profiling Workflow." *bioRxiv*. <https://doi.org/10.1101/2022.10.03.510579>.
- Portik, Daniel M, C Titus Brown, and N Tessa Pierce-Ward. 2022. "Evaluation of Taxonomic Classification and Profiling Methods for Long-Read Shotgun Metagenomic Sequencing Datasets." *BMC Bioinformatics* 23 (1): 541. <https://doi.org/10.1186/s12859-022-05103-0>.
- Quince, Christopher, Alan W Walker, Jared T Simpson, Nicholas J Loman, and Nicola Segata. 2017. "Shotgun Metagenomics, from Sampling to Analysis." *Nature Biotechnology* 35 (9): 833–44. <https://doi.org/10.1038/nbt.3935>.
- Rodriguez-R, Luis M, Santosh Gunturu, James M Tiedje, James R Cole, and Konstantinos T Konstantinidis. 2018. "Nonpareil 3: Fast Estimation of Metagenomic Coverage and Sequence Diversity." *mSystems* 3 (3). <https://doi.org/10.1128/mSystems.00039-18>.
- Rose, Rebecca, Olga Golosova, Dmitrii Sukhomlinov, Aleksey Tiunov, and Mattia Proserpi. 2019. "Flexible Design of Multiple Metagenomics Classification Pipelines with UGENE." *Bioinformatics (Oxford, England)* 35 (11): 1963–65. <https://doi.org/10.1093/bioinformatics/bty901>.
- Ruscheweyh, Hans-Joachim, Alessio Milanese, Lucas Paoli, Nicolai Karcher, Quentin Clayssen, Marisa Isabell Keller, Jakob Wirbel, et al. 2022. "Cultivation-Independent Genomes Greatly Expand Taxonomic-Profiling Capabilities of mOTUs Across Various Environments." *Microbiome* 10 (1): 212. <https://doi.org/10.1186/s40168-022-01410-z>.
- Schäffer, Alejandro A, Eric P Nawrocki, Yoon Choi, Paul A Kitts, Ilene Karsch-Mizrachi, and Richard McVeigh. 2018. "VecScreen_plus_taxonomy: Imposing a Tax(onomy) Increase on Vector Contamination Screening." *Bioinformatics (Oxford, England)* 34 (5): 755–59. <https://doi.org/10.1093/bioinformatics/btx669>.
- Schloss, Patrick D, Sarah L Westcott, Thomas Ryabin, Justine R Hall, Martin Hartmann, Emily B Hollister, Ryan A Lesniewski, et al. 2009. "Introducing Mothur: Open-Source, Platform-Independent, Community-Supported Software for Describing and Comparing Microbial Communities." *Applied and Environmental Microbiology* 75 (23): 7537–41. <https://doi.org/10.1128/AEM.01541-09>.
- Schmieder, Robert, and Robert Edwards. 2011. "Quality Control and Preprocessing of Metagenomic Datasets." *Bioinformatics (Oxford, England)* 27 (6): 863–64. <https://doi.org/10.1093/bioinformatics/btq111>.

746 [//doi.org/10.1093/bioinformatics/btr026](https://doi.org/10.1093/bioinformatics/btr026).

747 Schubert, Mikkel, Stinus Lindgreen, and Ludovic Orlando. 2016. "AdapterRemoval v2:
748 Rapid Adapter Trimming, Identification, and Read Merging." *BMC Research Notes*
749 9 (February): 88. <https://doi.org/10.1186/s13104-016-1900-2>.

750 Sczyrba, Alexander, Peter Hofmann, Peter Belmann, David Koslicki, Stefan Janssen,
751 Johannes Dröge, Ivan Gregor, et al. 2017. "Critical Assessment of Metagenome
752 Interpretation-a Benchmark of Metagenomics Software." *Nature Methods* 14 (11):
753 1063–71. <https://doi.org/10.1038/nmeth.4458>.

754 Sena Brandine, Guilherme de, and Andrew D Smith. 2021. "Falco: High-Speed FastQC
755 Emulation for Quality Control of Sequencing Data." *F1000Research* 8 (1874): 1874.
756 <https://doi.org/10.12688/f1000research.21142.2>.

757 Sharpton, Thomas J. 2014. "An Introduction to the Analysis of Shotgun Metagenomic
758 Data." *Frontiers in Plant Science* 5 (June): 209. <https://doi.org/10.3389/fpls.2014.00209>.

759 Shen, Wei, Hongyan Xiang, Tianquan Huang, Hui Tang, Mingli Peng, Dachuan Cai,
760 Peng Hu, and Hong Ren. 2023. "KMCP: Accurate Metagenomic Profiling of Both
761 Prokaryotic and Viral Populations by Pseudo-Mapping." *Bioinformatics* 39 (1):
762 btac845. <https://doi.org/10.1093/bioinformatics/btac845>.

763 Sim, Mikang, Jongin Lee, Daehwan Lee, Daehong Kwon, and Jaebum Kim. 2020.
764 "TAMA: Improved Metagenomic Sequence Classification Through Meta-Analysis."
765 *BMC Bioinformatics* 21 (1): 185. <https://doi.org/10.1186/s12859-020-3533-7>.

766 Straub, Daniel, Nia Blackwell, Adrian Langarica-Fuentes, Alexander Peltzer, Sven
767 Nahnsen, and Sara Kleindienst. 2020. "Interpretations of Environmental Micro-
768 bial Community Studies Are Biased by the Selected 16S rRNA (Gene) Amplicon
769 Sequencing Pipeline." *Frontiers in Microbiology* 11 (October): 550420. <https://doi.org/10.3389/fmicb.2020.550420>.

770 Sun, Zheng, Shi Huang, Meng Zhang, Qiyun Zhu, Niina Haiminen, Anna Paola Car-
771 rieri, Yoshiki Vázquez-Baeza, et al. 2021. "Challenges in Benchmarking Metage-
772 nomic Profilers." *Nature Methods* 18 (6): 618–26. <https://doi.org/10.1038/s41592-021-01141-3>.

773 Titus Brown, C, and Luiz Irber. 2016. "Sourmash: A Library for MinHash Sketching
774 of DNA." *Journal of Open Source Software* 1 (5): 27. <https://doi.org/10.21105/joss.00027>.

775 Uritskiy, Gherman V, Jocelyne DiRuggiero, and James Taylor. 2018. "MetaWRAP-a
776 Flexible Pipeline for Genome-Resolved Metagenomic Data Analysis." *Microbiome*
777 6 (1): 158. <https://doi.org/10.1186/s40168-018-0541-1>.

778 Vågene, Åshild J, Alexander Herbig, Michael G Campana, Nelly M Robles García,
779 Christina Warinner, Susanna Sabin, Maria A Spyrou, et al. 2018. "Salmonella
780 Enterica Genomes from Victims of a Major Sixteenth-Century Epidemic in Mex-
781 ico." *Nature Ecology & Evolution* 2 (3): 520–28. <https://doi.org/10.1038/s41559-017-0446-6>.

782 Veiga Leprevost, Felipe da, Björn A Grüning, Saulo Alves Aflitos, Hannes L
783 Röst, Julian Uszkoreit, Harald Barsnes, Marc Vaudel, et al. 2017. "Bio-
784 Containers: An Open-Source and Community-Driven Framework for Soft-
785 ware Standardization." *Bioinformatics (Oxford, England)* 33 (16): 2580–82.
786 <https://doi.org/10.1093/bioinformatics/btx192>.

- 792 Wick, Ryan R, Louise M Judd, Claire L Gorrie, and Kathryn E Holt. 2017. "Completing Bacterial Genome Assemblies with Multiplex MinION Sequencing." *Microbial Genomics* 3 (10): e000132. <https://doi.org/10.1099/mgen.0.000132>.
- 793
- 794
- 795 Wood, Derrick E, Jennifer Lu, and Ben Langmead. 2019. "Improved Metagenomic Analysis with Kraken 2." *Genome Biology* 20 (1): 257. <https://doi.org/10.1186/s13059-019-1891-0>.
- 796
- 797
- 798 Wratten, Laura, Andreas Wilm, and Jonathan Göke. 2021. "Reproducible, Scalable, and Shareable Analysis Pipelines with Bioinformatics Workflow Managers." *Nature Methods* 18 (10): 1161–68. <https://doi.org/10.1038/s41592-021-01254-9>.
- 799
- 800
- 801 Wright, Robyn J, André M Comeau, and Morgan G I Langille. 2023. "From Defaults to Databases: Parameter and Database Choice Dramatically Impact the Performance of Metagenomic Taxonomic Classification Tools." *Microbial Genomics* 9 (3). <https://doi.org/10.1099/mgen.0.000949>.
- 802
- 803
- 804
- 805 Ye, Simon H, Katherine J Siddle, Daniel J Park, and Pardis C Sabeti. 2019. "Benchmarking Metagenomics Tools for Taxonomic Classification." *Cell* 178 (4): 779–94. <https://doi.org/10.1016/j.cell.2019.07.010>.
- 806
- 807
- 808 Yilmaz, Pelin, Laura Wegener Parfrey, Pablo Yarza, Jan Gerken, Elmar Pruesse, Christian Quast, Timmy Schweer, Jörg Peplies, Wolfgang Ludwig, and Frank Oliver Glöckner. 2014. "The SILVA and 'All-Species Living Tree Project (LTP)' Taxonomic Frameworks." *Nucleic Acids Research* 42 (Database issue): D643–8. <https://doi.org/10.1093/nar/gkt1209>.
- 809
- 810
- 811
- 812