

nf-core/taxprofiler: highly parallelised and flexible pipeline for metagenomic taxonomic classification and profiling

Sofia Stamouli¹, Moritz E. Beber², Tanja Normark³, Thomas A. Christensen II⁴, Lili Andersson-Li⁵, Maxime Borry⁶, Mahwash Jamy⁷, nf-core community⁸, James A. Fellows Yates⁹

¹Department of Microbiology, Tumor and Cell Biology, Karolinska Institutet

¹Department of Clinical Microbiology, Karolinska University Hospital

²Unseen Bio ApS

³Department of Microbiology, Tumor and Cell Biology, Karolinska Institutet

³Department of Clinical Microbiology, Karolinska University Hospital

⁴Veterinary Diagnostic Laboratory, Kansas State University College of Veterinary Medicine

⁵Department of Microbiology, Tumor and Cell Biology, Karolinska Institutet

⁵Department of Clinical Microbiology, Karolinska University Hospital

⁶Department of Archaeogenetics, Max Planck Institute for Evolutionary Anthropology

⁷Department of Microbiology, Tumor and Cell Biology, Karolinska Institutet

⁷Department of Clinical Microbiology, Karolinska University Hospital

⁸

⁹Department of Archaeogenetics, Max Planck Institute for Evolutionary Anthropology

⁹Department of Paleobiotechnology, Leibniz Institute for Natural Product Research and Infection Biology Hans Knöll Institute (current address)

1 Abstract

Metagenomic classification tackles the problem of characterising the taxonomic source of all DNA sequencing reads in a sample. A common approach to address the differences and biases between the many different taxonomic classification tools is to run metagenomic data through multiple classification tools and databases. This, however, is a very time-consuming task when performed manually - particularly when combined with the appropriate preprocessing of sequencing reads before the classification.

Here we present nf-core/taxprofiler, a highly parallelised read-processing and taxonomic classification pipeline. It is designed for the automated and simultaneous classification and/or profiling of both short- and long-read metagenomic sequencing libraries against a 11 taxonomic classifiers and profilers as well as databases within a single pipeline run. Implemented in Nextflow and as part of the nf-core initiative, the pipeline benefits from high levels of scalability and portability, accommodating from

36 small to extremely large projects on a wide range of computing infrastructure. It has
37 been developed following best-practise software development practises and commu-
38 nity support to ensure longevity and adaptability of the pipeline, to help keep it up to
39 date with the field of metagenomics.

40 2 Introduction

41 Whole-genome, metagenomic sequencing offers strong benefits to the taxonomic clas-
42 sification of DNA samples over targeted approaches (Eloe-Fadrosh et al. 2016; Florian
43 P. Breitwieser, Lu, and Salzberg 2019). While metabarcoding approaches targeting
44 the 16S rRNA or other marker genes are widely used due to low cost and large, di-
45 verse reference databases (Yilmaz et al. 2014; Lynch and Neufeld 2015), metagenomic
46 approaches have been gaining popularity with the increasingly lower costs of, for
47 example, shotgun sequencing. These metagenomic analyses with whole microbial
48 genome as references have been shown to provide a similar level of taxonomic res-
49 olution (Hillmann et al. 2018). However they also have the added benefit of having
50 greater reusability potential of the data, such as for whole genome and/or functional
51 classification (Sharpton 2014; Quince et al. 2017).

52 Taxonomic classifiers (sometimes referred to as taxonomic bidders) aim to identify
53 the original ‘taxonomic source’ of a given DNA sequence (Ye et al. 2019; Meyer et al.
54 2022; Govender and Eyre 2022). In metagenomics, this typically consists of comparing
55 millions of DNA reads (sequenced DNA molecules) against hundreds or thousands
56 of reference genomes either via sequence alignment or ‘k-mer matching’ (Sharpton
57 2014; Sun et al. 2021). The reference genome with the most similar match to the
58 read is then considered the most likely original ‘source’ organism of that sequence. In
59 this article we will also refer to ‘taxonomic profilers’. We consider these as classifiers
60 that also try to infer sequence abundance (i.e. re-assignment of counts to the most
61 likely source based on the distribution of other hits) or biological relative abundance
62 of the organism in the original sample (by coverage of expected marker genes, copy
63 number estimations etc.), in addition to the simple read classification (Nayfach and
64 Pollard 2016). We will use classifiers and profilers interchangeably throughout the
65 publication.

66 Having to identify the original source of the many DNA sequences out of the many ref-
67 erence genomes in a time and computationally efficient manner is a difficult problem.
68 In many cases biologists are not just interested as to which organism of each DNA
69 sequence comes from, but also in using this information to infer the original ‘cellular’
70 (or natural) abundance of each organism of the given environment - something that
71 is very difficult due to the biases inherent to DNA extraction and sequencing. There-
72 fore a plethora of tools have been developed to address these challenges, all with their
73 own biases and specific contexts (Sczyrba et al. 2017; Meyer et al. 2022). Furthermore,
74 each tool often produces tool-specific output formats making it difficult to efficiently
75 cross compare results. Thus, no established ‘gold standard’ classifier tool or method
76 currently exists.

77 One solution to addressing the problem of choice among the range of different tools
78 is to run all of them in parallel, and cross compare the results. This can be useful both
79 for benchmarking studies (e.g. Sczyrba et al. 2017; Meyer et al. 2022), but also to
80 build consensus profiles whereby confidence of a particular taxonomic identification
81 can be increased when it is detected by multiple tools (McIntyre et al. 2017; Ye et al.
82 2019).

83 A second challenge in taxonomic classification (and arguably a larger one) is a ques-
84 tion of databases. As with tools, there is no one set ‘gold standard’ database. Different
85 questions and contexts require different databases, such as when a researcher wants
86 to search for both bacterial and viral species in samples, but as an extension of this,
87 taxonomic classifiers often will need different settings for each database. Further-
88 more, as genomic sequencing becomes cheaper and more efficient, the number of
89 publicly available reference genomes is rapidly increasing (Nasko et al. 2018). Conse-
90 quently, the size of reference databases of taxonomic classifiers is also growing, often
91 outpacing the computational capacity available to researchers. In fact, while this was
92 one of the main motivations behind classifiers such as Kraken2 (Wood, Lu, and Lang-
93 mead 2019), these algorithmic techniques are already becoming insufficient (Wright,
94 Comeau, and Langille 2023).

95 Finally, with the decrease of costs, the possibility for larger and larger metagenomic
96 sequencing datasets increases, leading to increasing sample sizes in studies. This is
97 exemplified by the doubling of the number of metagenomes on the European Bioin-
98 formatic Institute’s MGnify database within just two years (Mitchell et al. 2019).

99 Altogether this highlights the need for methods to efficiently profile many samples
100 using many tools. Manually setting up bioinformatic jobs for classification tasks for
101 each database and settings against different tools on traditional academic computing
102 infrastructure (e.g. high performance computing clusters or ‘HPC’ clusters) can be
103 very tedious. Additionally, particularly for very large sample sets, there is increas-
104 ing use of cloud platforms that have greater scalability than traditional HPCs. Being
105 able to reliably and reproducibly execute taxonomic classification tasks across infras-
106 tructure with minimal intervention would therefore be a boon for the metagenomics
107 field.

108 In recent years, workflow managers such as Nextflow (Di Tommaso et al. 2017) or
109 Snakemake (Mölder et al. 2021) have become highly popular in bioinformatics. These
110 frameworks provide for developers robust workflow execution with different HPC
111 scheduling tools and software provisioning systems, ensuring maximum portability
112 and efficient in different computational contexts. While a range of metagenomic
113 pipelines already exist (a non-exhaustive list being for example, StaG-mwc by
114 Boulund et al. 2023; MetaMeta by Piro, Matschkowski, and Renard 2017; TAMA by
115 Sim et al. 2020; UGENE by Rose et al. 2019; and Sunbeam by Clarke et al. 2019), few
116 leverage workflow managers to make multi-step workflows easier to use in HPC or
117 cloud infrastructure. Furthermore, often these pipelines aim to carry out multiple
118 different types of metagenomic analyses (e.g. also performing functional or assembly
119 analyses, such as Morais et al. 2022; Boulund et al. 2023) of which each step has
120 fewer options of tools and may execute functionality unwanted by the end user.

Here we present nf-core/taxprofiler (<https://nf-co.re/taxprofiler>), a pipeline designed to allow users to efficiently and simultaneously taxonomically classify or profile short- and long-read sequencing data. At the time of writing it supports 11 classifiers and an arbitrary number of databases per classifier in a single pipeline run. nf-core/taxprofiler utilises Nextflow (Di Tommaso et al. 2017) to ensure efficiency, portability, and scalability, and has been developed within the nf-core initiative of Nextflow pipelines (Ewels et al. 2020) to ensure high quality coding practises and user accessibility. It includes detailed documentation and a graphical-user-interface (GUI) execution interface in addition to a standard command-line-interface (CLI).

3 Description

nf-core/taxprofiler aims to facilitate three main steps of a typical whole-genome, metagenomic sequencing analysis workflow (Chiu and Miller 2019, Figure 1). A longer description of the available functionality and motivations can be seen in the [Supplementary Information](#).

In brief, nf-core/taxprofiler can accept short- (e.g. Illumina) and/or long-read (e.g. Nanopore) FASTQ or FASTA files. These are supplied to the pipeline in the form of a TSV file that includes basic sample and sequencing library metadata. The pipeline can then be executed either via a standard Nextflow command-line-interface execution or graphical-user-interface through either the open-source and free nf-core launch page (<https://nf-co.re/launch>) or the commercial (with free-tier) Nextflow tower (<https://tower.nf>) solution. Examples of the command-line execution and nf-core launch GUI can be seen in the [Supplementary Information](#).

The pipeline can perform a range of metagenomics appropriate read preprocessing steps, such as adapter removal, read merging, low-sequence complexity filtering, host- or contamination removal, and/or per-sample run merging. All of these steps are optional, and are aimed at removing possible sequencing artefacts that may result in false positive taxonomic classification hits or improve classification efficiency. Most of these steps also provide options of different tools to account for user preference.

After pre-processing, nf-core/taxprofiler can perform simultaneous profiling of pre-processed reads with up to as many as 11 different taxonomic classifiers or profilers (Table 1). Additionally on top of this, also simultaneously for each of the classifiers, an arbitrary number of databases as supplied by the user. As of version 1.1.0, the following classifiers and profilers are available: Kraken2 (Wood, Lu, and Langmead 2019), Bracken (Lu et al. 2017), KrakenUniq (F. P. Breitwieser, Baker, and Salzberg 2018), Centrifuge (Kim et al. 2016), MALT (Vågane et al. 2018), DIAMOND (Buchfink, Reuter, and Drost 2021), Kaiju (Menzel, Ng, and Krogh 2016), MetaPhlAn (Blanco-Míguez et al. 2023), mOTUs (Ruscheweyh et al. 2022), ganon (Piro et al. 2020), and KMCP (Shen et al. 2023). Databases are also supplied via a input TSV file, which also allows per-database custom classification parameters - meaning a given database can be supplied multiple times each with different parameters or multiple different databases per profiler. All classifiers with secondary steps to generate or convert to

162 additional output file formats are also included.

163 Post-processing of taxonomic profiles include standardisation and aggregation of pro-
 164 files , i.e. merging of multiple profiles into a single multi-sample table for easier com-
 165 parison between profilers, with the tool TAXPASTA (Beber et al. 2023), and visualisa-
 166 tion of profiles with Krona (Ondov, Bergman, and Phillippy 2011) where supported.

167 All relevant preprocessing statistics are displayed in an interactive and dynamic Mul-
 168 tiQC report (Ewels et al. 2020).

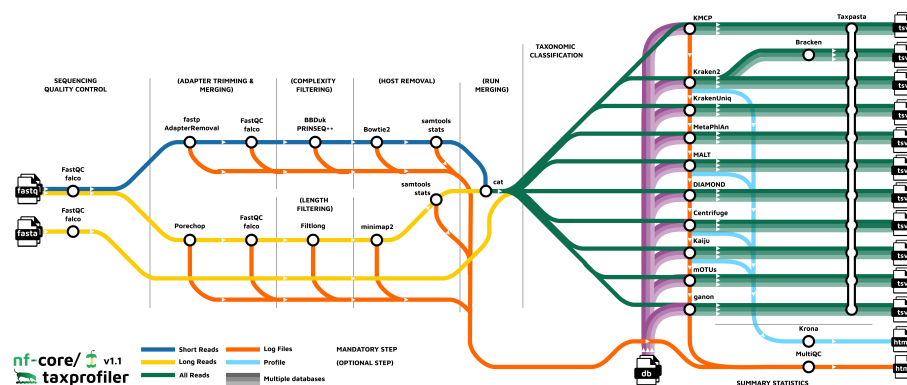


Figure 1: Visual overview of the nf-core/taxprofiler workflow. nf-core/taxprofiler can take in FASTQ (short or long reads) or FASTA files (long reads), that will optionally go through sequencing quality control (e.g. with FastQC), read preprocessing (e.g. removal of adapters), complexity filtering, host removal, and run merging before performing taxonomic classification and/or profiling with a user-selected range of tools and databases. Output from all classifiers and profilers are standardised into a common taxon table format, and when supported visualisations of the profiles are generated.

Table 1: List of nf-core/taxprofiler supported taxonomic/classifiers profilers as of version 1.1 and their approximate method and supported input database types. Primary algorithm refers to the algorithm type used for sequencing matching. Reference type refers to the typical sequence type used in database construction of the tool. Sequencing matching type refers to which ‘molecular alphabet’ is primarily used for matching between a query (read) and a reference (genome/gene).

Tool	Primary Algorithm	Reference Type	Sequence Matching Type
Kraken2	k-mer based	whole-genome	Nucleotide
Kaiju	k-mer based	whole-genome	Amino Acid
Bracken	k-mer based	whole-genome	Nucleotide
KrakenUniq	k-mer based	whole-genome	Nucleotide
ganon	k-mer based	whole-genome	Nucleotide
KMCP	k-mer based	whole-genome	Nucleotide

Tool	Primary Algorithm	Reference Type	Sequence Matching Type
MALT	alignment based	whole-genome	Nucleotide/Amino Acid
DIAMOND	alignment based	whole-genome	Amino Acid
Centrifuge	alignment based	whole-genome	Nucleotide
MetaPhlAn	alignment based	marker-gene	Nucleotide
mOTUS	alignment based	marker-gene	Nucleotide

169 nf-core/taxprofiler comes with extensive documentation for general usage, short- and
170 long- parameter help texts, and output file descriptions. To ensure maximum accessi-
171 bility, these are available in pipeline results as markdown files ([https://github.com/nf-](https://github.com/nf-core/taxprofiler)
172 [core/taxprofiler](https://github.com/nf-core/taxprofiler)), on the nf-core website (<https://nf-co.re/taxprofiler>) and for the pa-
173 rameter help texts on the command line via standard --help. The output documen-
174 tation also aims to guide users as the most suitable files for different types of down-
175 stream analysis

176 4 Discussion

177 A range of pipelines already exists for taxonomic profiling, however, each have
178 their own particular purpose and capabilities. We compared the functionality
179 of nf-core/taxprofiler against four other recently published or released pipelines,
180 selected based on their similarity of purpose to nf-core/taxprofiler. The selection
181 criteria and a more detailed comparison between the five pipelines can be seen
182 in the [Supplementary Information](#). Overall, while there was a general similarity
183 across all pipelines, nf-core/taxprofiler showed the largest number of options for
184 pipeline execution accessibility, and user choice. This is facilitated through the
185 use of an established workflow manager (with Nextflow supporting 7 software
186 environment/container systems), support for both CLI and GUI execution, and by the
187 number of supported classifiers. Furthermore, it is unique in that is the only pipeline
188 to support supplying multiple database for all of the tools in a single pipeline run.

Table 2: Comparison of functionality with four recent taxonomic pipelines with simi-
lar functionality. A more detailed textual comparison can be found in the [Supplemen-
tary Information](#). Category keys are as follows: I - Information, R - Reproducibility,
A - Accessibility, P - Portability, S - Scalability, F - Functionality.

Category	Criterion	StaG- mwc	sunbeam	Unipro UGENE	tama	nf- core/taxprofiler
I	Source code URL	https://github.com/ctmrbio/stag-mwc	https://github.com/sunbeam-labs/sunbeam	https://github.com/ugeneuniprgkimlab/ugene	https://github.com/TAMA	https://github.com/nf-core/taxprofiler/

Category	Criterion	StaG-mwc	sunbeam	Unipro UGENE	tama	nf-core/taxprofiler
I	Evaluated version	0.7.0	4	48	githash: 3a22c8f	1.1.0
I	Last release date	2023-06-13	2023-08-08	2023-08-08	2022-03-02	2023-09-19
I	Publication year	Unpublished	2019	2019	2020	This publication
I	Publication DOI	Unpublished	10.1186/s40168-019-0658-x	10.1093/bioinformatics/btz259	10.1186/s13052-020-3533-7	This publication
R	Pipeline versioning	Yes	Yes	Yes	No	Yes
R	Software versioning	Yes	Yes	Yes	Yes	Yes
R	Nr. software environments or container engines supported	2	2	0	1	7
A	Installation documentation	Yes	Yes	Yes	Yes	Yes
A	Usage documentation	Yes	Yes	Yes	Yes	Yes
A	Output documentation	Yes	Yes	Yes	Yes	Yes
A	CLI execution interface	Yes	Yes	No	Yes	Yes
A	GUI execution interface	No	No	Yes	No	Yes
A/S	Integration a scheduling systems	Yes	Yes	No	No	Yes
P/A	Nr. supported operating systems	2	1	3	1	2
P	Local machine integration	Yes	Yes	Yes	Yes	Yes
P/S	HPC scheduler integration	Yes	Yes	No	No	Yes

Category	Criterion	StaG-mwc	sunbeam	Unipro UGENE	tama	nf-core/taxprofiler
P/S	Cloud computing integration	Unsure	Unsure	No	No	Yes
P/S	Integration with multiple scheduling systems	Partial	Partial	No	No	Yes
S	Per-process resource optimisation	Yes	Yes	Yes	No	Yes
F	Short read support	Yes	Yes	Yes	Yes	Yes
F	Long read support	No	No	Yes	No	Yes
F	Read preprocessing	Yes	Yes	Yes	Yes	Yes
F	Sequencing depth estimation	Yes	No	No	No	No
F	Complexity filtering	No	Yes	No	No	Yes
F	Host removal	Yes	Yes	Partial	No	Yes
F	Nr. supported taxonomic classifiers/profilers	7	3	3	3	11
F	Graphical run reports	Yes	No	No	No	Yes
F	Standardised profiles	No	No	No	Yes	Yes
F	Multiple database supported	Partial	No	No	No	Yes
F	Metagenomic assembly	No	Yes	No	No	No
F	Visualisation	No	No	No	No	Partial

189 Another important advantage of nf-core/taxprofiler is that it is being developed
 190 within the nf-core community (<https://nf-co.re>), that provides strong long-term
 191 support for the continued community-based development and maintenance of its
 192 pipelines. In this framework, we will continue to add additional preprocessing,
 193 metagenomic classification, and profiling tools as they become established and as

194 requested by the metagenomics community. For example, we feel that the inclusion
195 of steps such as sequencing saturation estimation as already being performed
196 by a similar pipeline StaG-mwc (<https://github.com/ctmrbio/stag-mwc>) would be
197 beneficial to the nf-core/taxprofiler workflow (possibly with dedicated tools such as
198 Nonpareil, Rodriguez-R et al. 2018), and/or more performant complexity filtering
199 tools such as Komplexity as offered by the sunbeam metagenomics pipeline (Clarke
200 et al. 2019). Additional tools that could be added for short-read classification could
201 include sourmash (Titus Brown and Irber 2016) that provides scalable sequence
202 to sequence comparison or other marker gene reference tools such as tools such
203 as METAXA2 (Bengtsson-Palme et al. 2015) that use shotgun sequencing reads to
204 recover 16S sequences from metagenomic samples. Adding additional classifiers also
205 applies to extend support to other sequencing platforms; nf-core/taxprofiler already
206 supports Nanopore long-read data, however the use of long-read PacBio data for
207 metagenomic data is growing in interest (Portik, Brown, and Pierce-Ward 2022).
208 We are therefore considering adding dedicated preprocessing steps for this type of
209 sequencing data.

210 A remaining major challenge for metagenomics researchers (and not supported in
211 the same workflow by any of the compared pipelines above) is the construction of
212 databases for each profiling tool. Given there still are no curated, high-quality ‘gold
213 standard’ databases in metagenomics, and while nf-core/taxprofiler allows the pro-
214 filing against multiple databases and settings in parallel, currently the pipeline still
215 requires users to construct these manually and to supply to the pipeline. While we
216 feel this is currently a reasonable investment as such databases are typically repeat-
217 edly re-used, we are exploring the possibility to add an additional complementary
218 workflow in the pipeline to allow automated database construction of all classifica-
219 tion tools, given a set of FASTA reference files.

220 Finally, once an overall taxonomic profile is generated, researchers often wish to val-
221 idate hits through more sensitive and accurate methods such as with read-mapping
222 alignment. While read alignment is supported by other pipelines such as StaG-mwc,
223 this happens in-parallel to the taxonomic profiling and requires prior expectation of
224 which reference genomes to map against. Instead, nf-core/taxprofiler could be eas-
225 ily extended to have a validation step similar to the approach of the ancient DNA
226 metagenomic pipeline aMeta (Pochon et al. 2022). Utilising Nextflow’s execution par-
227 allelism, the input sequences could be aligned back to the reference genomes of only
228 those species with hits resulting from the taxonomic classification, but with dedicated
229 accurate short- or long-read aligners. In addition to the more precise classification,
230 post-classification read-alignment could also be particularly useful for researchers in
231 palaeogenomics who wish to use tools other than KrakenUniq for initial classification
232 (as in aMeta), where alignment information can be used to authenticate ancient DNA
233 within their samples, but also in clinical metagenomics to identify potential pathogens
234 at much finer resolution (e.g. down to strain level).

235 Another motivation for developing nf-core/taxprofiler, despite the large number of ex-
236 isting metagenomics pipelines, is that by establishing a taxonomic profiling pipeline
237 within the nf-core ecosystem, it is possible to begin building both standalone but

238 also an integrated suite of powerful interconnected pipelines for the major stages
239 of metagenomic workflows. Existing microbial- and metagenomics- related pipelines
240 within the nf-core initiative include nf-core/ampliseq (Straub et al. 2020), nf-core/mag
241 (Krakau et al. 2022), and nf-core/funcscan (<https://nf-co.re/funcscan>). We expect over
242 time the ability to link inputs and outputs of each workflow to develop comprehensive
243 metagenomic analyses, while still maintaining powerful standalone pipelines, provid-
244 ing maximal user choice but with familiar interfaces.

245 5 Conclusion

246 nf-core/taxprofiler is an accessible, efficient, and scalable pipeline for metagenomic
247 taxonomic classification and profiling that can be executed on anywhere from laptops
248 to the cloud. To our knowledge, the pipeline offers the largest number of taxonomic
249 profilers across similar pipelines, providing flexibility for users not just on choice of
250 profiling tool but also with databases and database settings within a single run. With
251 the development within the open and welcoming nf-core community and with best-
252 practise development infrastructure, we look forward to further contributions and in-
253 volvement of the wider metagenomics community, and also we hope that through de-
254 tailed documentation and a range of execution options, nf-core/taxprofiler will make
255 reproducible and high-throughput metagenomics more accessible for a wide range of
256 disciplines.

257 6 Code Availability

258 nf-core/taxprofiler source code is available on GitHub at [https://github.com/nf-core/](https://github.com/nf-core/taxprofiler)
259 [taxprofiler](https://github.com/nf-core/taxprofiler), and each release is archived on Zenodo (latest version DOI: [10.5281/zen-](https://doi.org/10.5281/zenodo.7728364)
260 [odo.7728364](https://doi.org/10.5281/zenodo.7728364))

261 The version of the pipeline described in this paper is version 1.1.0 (release specific
262 Zenodo archive DOI: [10.5281/zenodo.8358147](https://doi.org/10.5281/zenodo.8358147))

263 7 Acknowledgments

264 We thank Prof. Christina Warinner and the Microbiome Sciences group MPI-EVA for
265 original discussions that lead to the pipeline. We are also grateful for the nf-core
266 community for the original and ongoing support in the development in the pipeline, in
267 particular for the contributions by Lauri Mesilaakso, Jianhong Ou, and Rafał Stepień.

268 8 Funding

269 S.S. and L.A-L. were supported by Rapid establishment of comprehensive laboratory
270 pandemic preparedness – RAPID-SEQ. This material is based upon work supported by
271 the U.S. Department of Agriculture, Agricultural Research Service, under agreement

272 No. 58-3022-0-001 (T.A.C II). M.B. and J.A.F.Y were supported by the Max Planck So-
273 ciety. J.A.F.Y was supported by the Werner Siemens-Stiftung ("Paleobiotechnology",
274 Awarded to Prof. Pierre Stallforth and Prof. Christina Warinner).

275 9 Supplementary Information

276 9.1 Implementation

277 9.1.1 Input and Execution

278 The pipeline can be executed via typical Nextflow commands (Code Block 1), or us-
279 ing the standard nf-core 'launch' GUI (Figure 2), making the pipeline accessible for
280 both computationally experienced as well as less experienced researchers. In addi-
281 tion to the general usage and parameter documentation of the pipeline ([https://nf-](https://nf-co.re/taxprofiler)
282 [co.re/taxprofiler](https://nf-co.re/taxprofiler)). The GUI offers immediate assistance and guidance to users on
283 what each parameter does, both in short- and long-form, with long-form parameter
284 descriptions additionally describing which tool-specific parameters are being modi-
285 fied for each pipeline parameter (<https://nf-co.re/launch/?pipeline=taxprofiler>). The
286 GUI also includes controlled user input by providing strict drop-down lists and input
287 validation prior execution of the pipeline (Figure 2) to reduce the risk of typos and
288 other mistakes, which is in contrast to the command-line interface that only includes
289 validation at pipeline run-time.

Listing 1 Example nf-core/taxprofiler command for running short-read quality control, removal of host DNA and executing the k-mer based Kraken2 and marker gene alignment MetaPhlAn tools.

```
$ nextflow run nf-core/taxprofiler \  
-r 1.1.0 \  
-profile singularity,<institute> \  
--input <samplesheet.csv> \  
--databases <database.csv> \  
--perform_shortread_qc \  
--shortread_qc_minlength 20 \  
--preprocessing_qc_tool falco \  
--run_host_removal --hostremoval_reference 'host_genome.fasta' \  
--run_kraken2 --kraken2_save_reads \  
--run_metaphlan \  
--run_krona \  
--run_profile_standardisation
```

290 An example nf-core command line execution of the pipeline can be seen in Code
291 Block 1, where two input files are supplied: one file specifying paths of FASTQ files
292 of metagenomic samples and necessary metadata for preprocessing (such as sample

Preprocessing short-read QC options

Launch

--shortread_qc_minlength

15

?

Specify the minimum length of reads to be retained

Specifying a minimum read length filtering can speed up profiling by reducing the number of short unspecific reads that need to be match/aligned to the database.

Modifies tool parameter(s):

- removed from reads --length_required
- AdapterRemoval: --min length

--perform_shortread_complexityfilter

☐ True
☒ False

?

Turns on nucleotide sequence complexity filtering

--shortread_complexityfilter_tool

bbduk

?

Specify which tool to use for complexity filtering

[Select an option]
bbduk
prinseqplusplus
fastp

--shortread_complexityfilter_entropy

?

Specify the minimum sequence entropy level for complexity filtering

--shortread_complexityfilter_bbduk_windowsize

50

?

On this page

Nextflow command-line flags

> Input/output options

Preprocessing general QC options

Preprocessing short-read QC options

Preprocessing long-read QC options

Preprocessing host removal options

Preprocessing run merging options

Profiling options

Postprocessing and visualisation options

Show hidden params

Figure 2: Screenshot of the nf-core pipeline launch graphical user interface with nf-core/taxprofiler options displayed. The web browser-based interface provides guidance for how to configure each pipeline parameter by providing both short and long help descriptions to help guide users in which contexts to configure each parameter. Additional elements such as radio buttons, drop down menus, and background regular expressions check for validity of input. When pressing launch, a prepared configuration file and command is provided that can be copied and pasted by the user into the terminal

293 ID and sequencing platform), and the second file specifying paths to the user-defined
 294 databases with per-database classification parameters. Various parameters are avail-
 295 able to select different preprocessing steps, and provide additional configuration such
 296 as tool selection and value options. Note that even if a user supplies a given database
 297 in the database input sheet, the corresponding profiling tool must still be activated
 298 with the corresponding pipeline parameter (e.g. `--run_kraken2`). Per-classifier flags
 299 are also available for the optional saving of additional non-profile output files. Alter-
 300 natively to command line flags, parameters can be specified via pre-configured YAML
 301 format files, with which (provided no hardcoded paths are included) can be re-used
 302 across pipeline runs.

303 All nf-core pipelines are strictly versioned (specified with the Nextflow `-r` flag), and to
 304 ensure reproducibility, each version of the pipeline has a fixed set of software used for
 305 each step of the pipeline. The fixed set of software are controlled through the use of
 306 the conda package manager or containers (Docker, or Apptainer - previously known
 307 as Singularity, etc) from the stable Bioconda (Grüning et al. 2018) or BioContainers
 308 (Veiga Leprevost et al. 2017) repositories. This, coupled with the intrinsic Nextflow
 309 ability to execute on most infrastructure whether that is a local laptop (resource re-
 310 quirements permitting), traditional HPC, as well across common cloud providers also
 311 makes nf-core/taxprofiler a very portable pipeline that can be used in many contexts.

312 9.1.2 Preprocessing

313 Preprocessing steps in nf-core/taxprofiler are aimed at removing laboratory and se-
 314 quencing artefacts that may influence taxonomic profiling, either for computing re-
 315 source consumption or and/or false-positive or false-negative classification reasons.
 316 First sequencing quality control with FastQC (Andrews 2010) or Falco (Sena Brandine
 317 and Smith 2021) is carried out. Falco was included for reduced memory requirements,
 318 in particular for long read sequencing data. Artificial library adapter sequences added
 319 during sequencing reduce sequencing matching accuracy by reducing sequence speci-
 320 ficity, and in some cases, may result in false-positive hits due to adapter sequence con-
 321 tamination in reference genomes (Schäffer et al. 2018; F. P. Breitwieser, Baker, and
 322 Salzberg 2018) ¹. Additionally, paired-end merging may provide longer sequences
 323 that will allow for more specific classification when paired-end alignment is not sup-
 324 ported by a given classifier. For these tasks nf-core/taxprofiler can apply either fastp
 325 (Chen et al. 2018) or AdapterRemoval2 (Schubert, Lindgreen, and Orlando 2016) for
 326 short reads, and currently Porechop (Wick et al. 2017) for Oxford Nanopore long-read
 327 data. For both short and long reads, FastQC or Falco is run again to allow assessment
 328 on the performance of the adapter removal and/or pair-merging step.

329 Low complexity sequences, e.g. sequences containing long stretches of mono- or

¹For an ‘infamous’ case of adapter sequences in a published eukaryotic genome, see the following blog posts

Graham Etherington: <https://web.archive.org/web/20201219022000/http://grahametherington.blogspot.com/2014/09/why-you-should-qc-your-reads-and-your.html?m=1> why-you-should-qc-your-reads-and-your.html
 Sixing Huang: <https://web.archive.org/web/20220904205331/https://dgg32.medium.com/carp-in-the-soil-1168818d2191>
 (Accessed 2023-08-25)

330 di-nucleotide repeats provide little specific genetic information that contribute to
 331 taxonomic identification, as they can align to many different reference genomes
 332 (Schmieder and Edwards 2011; Clarke et al. 2019). Including such reads during
 333 taxonomic profiling can increase run-time and memory usage for little gain, as
 334 during lowest-common-ancestor (LCA) classification steps they will be assigned to
 335 high-level taxonomic ranks (e.g. Kingdom). `nf-core/taxprofiler` performs removal of
 336 these reads through complexity filtering algorithms as provided by `fastp`, `BBDuk`
 337 (Bushnell 2022), or `PRINSEQ++` (Cantu, Sadural, and Edwards 2019). Long read
 338 sequences often do not have such reads, as lengths are sufficient enough to capture
 339 greater sequence diversity - but it is sometimes desirable to only classify reads longer
 340 than a certain length - as these provide more precise taxonomic information (Dilthey
 341 et al. 2019; Portik, Brown, and Pierce-Ward 2022). Therefore, `nf-core/taxprofiler` can
 342 remove reads shorter than a user-defined length using `Filtlong`.

343 Removing host DNA is another common preprocessing step in metagenomic studies.
 344 This can help speed up run-time, particularly in microbiome studies, where detection
 345 of microbes are of interest. Furthermore, host-contamination of reference genomes in
 346 public databases is common (Longo, O'Neill, and O'Neill 2011; Kryukov and Imanishi
 347 2016; Florian P. Breitwieser et al. 2019). Therefore, the removal of such sequences can
 348 help decrease the risk of false positive taxonomic assignment. To remove multiple
 349 hosts or other sequences, all reference genomes can be combined into a single FASTA
 350 reference file. Short read host removal can be carried out with `Bowtie2` (Langmead
 351 and Salzberg 2012; Langmead et al. 2019) and `minimap2` (Li 2018) for long reads, both
 352 in combination with `SAMtools` (Li et al. 2009; Danecek et al. 2021), where reads are
 353 aligned against the reference genome and the off-target (unaligned) reads are then
 354 converted back to FASTQ format for classification.

355 Finally, `nf-core/taxprofiler` can optionally perform 'run merging' where multiple
 356 FASTQ files from the same sample but have been sequenced over multiple lanes are
 357 concatenated together to generate one profile per sample or library. The final set of
 358 reads used for profiling can be optionally saved for downstream re-use. Throughout
 359 all steps, relevant statistics and log files are generated and used both for the final
 360 pipeline run report as well as saved into the results directory of the pipeline run for
 361 further inspection where necessary.

362 9.1.3 Profiling

363 There are many types of metagenomic profiling techniques, from profiling against
 364 whole-genome references with alignment or k-mer based approaches, to methods in-
 365 volving alignment to species-specific marker-gene families (Quince et al. 2017; Ye et
 366 al. 2019). `nf-core/taxprofiler` aims to support and include all established classification
 367 or profiling tools as requested by the community.

368 The choice of tools used in a pipeline run is up to the user, with a tool being executed
 369 when both the corresponding database and `--run_<tool>` flag is provided. Specific
 370 classification settings for each tool and database are specified in the database CSV
 371 input sheet. Some tools also have pipeline level command-line flags for controlling

372 certain aspects of output files.

373 The following classifiers and profilers are supported in version 1.1.0 of nf-
374 core/taxprofiler: Kraken2 (Wood, Lu, and Langmead 2019), Bracken (Lu et al.
375 2017), KrakenUniq (F. P. Breitwieser, Baker, and Salzberg 2018), Centrifuge (Kim et
376 al. 2016), MALT (Vågane et al. 2018), DIAMOND (Buchfink, Reuter, and Drost 2021),
377 Kaiju (Menzel, Ng, and Krogh 2016), MetaPhlAn (Blanco-Míguez et al. 2023), mOTUs
378 (Ruscheweyh et al. 2022), ganon (Piro et al. 2020), KMCP (Shen et al. 2023).

379 By default, nf-core/taxprofiler produces the default per-sample taxonomic classifica-
380 tion profile output from a tool or a tool’s report generation tool. The output is nor-
381 mally in the form of counts per reference sequencing, with additional statistics about
382 the hits of a particular organism (estimated sequence abundance, taxonomic level etc.).
383 Users can also optionally request output of per-read classification output and output
384 such as classified and unclassified reads in FASTQ format, where supported.

385 The pipeline provides high efficiency, particularly during the metagenomic classifica-
386 tion stage, through the inherent parallelisation provided by Nextflow. While metage-
387 nomic classification is comparatively computationally intensive (in terms of mem-
388 ory and execution time; due to a combination of sequencing depth and number of
389 reference genomes), Nextflow automatically optimises the execution order of all the
390 steps in pipeline, maximising the number parallel running of multiple profilers and/or
391 databases at any given time point, as far as the available computational resources al-
392 low. For local machines such as laptops or desktops, Nextflow will automatically
393 detect all available computational resources, but this is customisable using Nextflow
394 configuration files. For HPC and cloud infrastructure, users typically have to define
395 the computational infrastructural environment the pipeline is being executed on (CPU
396 or memory limitations, queues, instance types, etc.). To facilitate the pipeline compu-
397 tational configuration, nf-core/taxprofiler supports use of more than 90 pre-defined
398 centralised generic and pipeline-specific institutional Nextflow configurations as pro-
399 vided by nf-core/configs (<https://nf-co.re/configs>). However, of course users are still
400 welcome to supply their own custom configuration files as with any typical Nextflow
401 run, further refining computational limitations or execution specifications.

402 9.1.4 Post-profiling

403 In metagenomic studies, it is common practise to compare the profiles among many
404 samples, and the results of multiple profiles are normally stored in ‘taxon tables’, i.e,
405 counts per reference taxon (rows), for each sample (columns). When available, nf-
406 core/taxprofiler supports the option to produce the ‘native’ taxon table of each classi-
407 fication tool when multiple samples are run.

408 One of the challenges that researchers face when comparing multiple taxonomic clas-
409 sifiers or profilers is the heterogenous output formats that are produced, that often
410 require custom parsing and merging scripts for each tool to standardise. To facilitate
411 more user-friendly cross-comparisons between tools, nf-core/taxprofiler utilises the
412 TAXPASTA tool (Beber et al. 2023) to generate standardised profiles and generate
413 multi-sample tables.

Summary statistics for the entire pipeline are visualised and displayed in a customisable MultiQC report (Ewels et al. 2020). When supported, quality control of data and pipeline runs are shown for manual verification. Krona plots (Ondov, Bergman, and Phillippy 2011) can also optionally be generated for supported tools to help provide further visualisation of taxonomic profiles.

9.1.5 Output

To summarise, the main default output from nf-core/taxprofiler are both classifier ‘native’ and standardised single- and multi-sample taxonomic profiles with counts per-taxon and an interactive MultiQC run report with all run statistics, in addition to the raw log files themselves where available.

The MultiQC run report displays statistics and summary visualisations for all steps of the pipeline where possible, lists of versions for all tools of each step of the pipeline. It also provides a dynamically-constructed text for the recommended ‘methods’ for reporting how the pipeline was executed (including relevant citations) that users can use in their own publications.

Optional outputs can include other types of profiles (e.g. per read classification) and in other formats as produced by the tools themselves, as well as raw reads from pre-processing steps and output visualisations from Krona. Nextflow resource usage and trace reports are also by default produced for users to check pipeline performance.

9.2 Comparison with other solutions

nf-core/taxprofiler has been specifically developed for the analysis of whole-genome, *metagenomic* sequencing data. While other types of taxonomic profiling data such as 16S amplicon sequencing are well established fields with a range of popular high-quality and best-practise tools pipelines (e.g. Blanco-Míguez et al. 2023; Schloss et al. 2009) and databases (DeSantis et al. 2006; Yilmaz et al. 2014), ‘gold standard’ tools and databases for metagenomics remain much less established. Thus, the need for highly-multiplexed classification is more desirable for the newer metagenomics methods.

We searched Google Scholar for open-source pipelines published or released in the last 5 years (at the time of writing, since 2018) that were designed primarily for metagenomic classification screening, that supported at least 2 classifiers, had at least one preprocessing step and were not specifically targeted at read classification of specific domains of taxa (e.g. viruses or bacteriophages only). We also included an additional open-source but unpublished pipeline at the recommendations of the authors of the pipeline due to the functional overlap to nf-core/taxprofiler. We then evaluated the pipelines based on their publications and documentation for typical metagenomic profiling workflow steps. We used a range of criteria related to expectations of modern bioinformatic workflows that can be summarised in the following four categories: reproducibility, accessibility, scalability, and portability (Wratten, Wilm, and Göke 2021). After searching, we selected the following pipelines for comparison with nf-core/taxprofiler that matched the specific criteria described above: sunbeam (v4,

Clarke et al. 2019), Unipro UGENE (v48, Rose et al. 2019), TAMA (githash: 3a22c8f, Sim et al. 2020), and StaG-mwc (0.7.0, Boulund et al. 2023).

In terms of accessibility, all pipelines have documentation describing the installation steps, usage instructions, and output files. However, there are varying levels of detail and comprehensiveness. In particular, StaG-mwc and nf-core/taxprofiler have the most detailed descriptions of all possible output files for every supported module, whereas Unipro UGENE and sunbeam have very minimal to possibly unfinished output documentation. For execution options, most of the pipelines provide CLI execution, except for Unipro UGENE which offers only GUI-based pipeline set-up (despite a command-line execution of the GUI generated configuration). In particular, nf-core/taxprofiler is the only pipeline providing both CLI and GUI interfaces for pipeline run execution.

Criteria covering portability also overlap with accessibility, as it implies options for and ease of different users running on different types of computing infrastructure, whether that is on their own laptop, on an HPC cluster, or in the cloud. Unipro UGENE is the only pipeline that explicitly states support for execution on all three major operating systems (Linux, OSX, Windows), whereas StaG-mwc and nf-core/taxprofiler can be run on unix operating systems (albiet possibly on Windows via Windows Subsystem for Linux (WSL)), and sunbeam and TAMA are only being supported on Linux.

While all pipelines support 'local' machine execution (e.g. personal laptops or desktops), a large portion of academic users execute computationally intensive bioinformatic tasks on HPC clusters. In these contexts, pipeline task submissions are normally managed by job schedulers, thus integration with schedulers is an important criterion for running large multi-step and parallelised pipelines. The three pipelines leveraging workflow managers (Snakemake and Nextflow) support integration with schedulers (StaG-mwc, sunbeam, and nf-core/taxprofiler) with nf-core/taxprofiler supporting the most by far (>10 [scheduling systems](#)) as natively offered by Nextflow. This allows the greatest possible choice for users in terms of which HPC infrastructure they can execute their pipeline on. As an extension of this, only nf-core/taxprofiler has explicit support for cloud computing (e.g. AWS, GCP, or Microsoft Azure) as provided by Nextflow, again maximising user choice and portability when it comes to running the pipeline.

In terms of scalability, the aforementioned integration with schedulers and cloud computing support implicitly maximises efficiency and parallelisation of pipeline runs, providing good scalability for varying numbers of input files and steps in the pipeline. Again, the three workflow manager based pipelines provide scalability, whereas there is no mention neither Unipro UGENE nor TAMA in reference to parallel task execution. Furthermore, all pipelines except TAMA, allowed per-process customisation of computational resources, something critical for maximising efficient scalability to ensure only the necessary resources for a given step of a pipeline are requested.

In terms of reproducibility, all five pipelines are good at ensuring reproducibility in terms of pipeline and software versioning (allowing re-execution of pipeline runs using the same software), with only TAMA not having stable versioned releases. How-

498 ever, installing software manually across different infrastructures can result in vari-
499 ability in the execution of each software² (Di Tommaso et al. 2017). The current most
500 popular solution to the problem of inconsistent software environments is to use con-
501 tainer engines such as Docker or Apptainer to run container images which are iso-
502 lated, deterministic computing environments which can be executed by any system
503 providing a container runtime. Only Unipro UGENE does not document the use of a
504 container system, with nf-core/taxprofiler offering the biggest choice for users, again,
505 courtesy of Nextflow with 6 different engine systems at the time of writing.

506 Finally, we compared metagenomics related functionality between the pipelines. All
507 pipelines support short-read FASTQ input, but only nf-core/taxprofiler explicitly re-
508 ports long-read support, while the documentation in Unipro UGENE states that assem-
509 bled contigs are possible input to some of the profilers. All pipelines support read pre-
510 processing (adapter clipping, and merging). In terms of tools used for preprocessing,
511 Trimmomatic (Bolger, Lohse, and Usadel 2014) is popular across the other pipelines
512 but is not supported in nf-core/taxprofiler. Only sunbeam and nf-core/taxprofiler sup-
513 port complexity filtering to remove low sequence diversity reads. In fact within sun-
514 beam, the authors developed their own dedicated, performant complexity filtering
515 tool Komplexity (Clarke et al. 2019). Most pipelines support some form of host re-
516 moval (only TAMA did not support this), and it is likely possible with Unipro UGENE
517 (although not directly described). In all cases, host removal consists of mapping pro-
518 cessed reads with an aligner and using the off-target reads for downstream profiling
519 (as implemented in nf-core/taxprofiler), however StaG-mwc has an additional sepa-
520 rate metagenomic host removal step with Kraken2. nf-core/taxprofiler supports by
521 far the largest number of taxonomic classifiers and profilers at 11 as of v1.1.0 - pro-
522 viding the greatest choice to users - with StaG-mwc offering 7, and the remaining
523 pipelines only 3. Only nf-core/taxprofiler and partly StaG-mwc explicitly support run-
524 ning each profiler with multiple databases. nf-core/taxprofiler is the only pipeline that
525 supports running an arbitrary number of different metagenomic profiler databases
526 each with their own settings. This makes it a useful for tool parameter compari-
527 son, testing different databases, or reducing the size of each database (e.g. per do-
528 main) to make it more flexibility for running on smaller computational infrastructure.
529 StaG-mwc allows multiple references for their short-read alignment steps rather than
530 the metagenomic profilers. For output, nf-core/taxprofiler, StaG-mwc, and sunbeam
531 (via an extension) support a singular run report for summarising all preprocessing
532 step. Only nf-core/taxprofiler and TAMA produce standardised output for all taxo-
533 nomic profilers, the former with the dedicated standalone tool TAXPASTA (Beber et
534 al. 2023). However Unipro UGENE additionally offers a ‘consensus’ profile using
535 WEVOTE (Metwally et al. 2016).

536 To summarise, many of the pipelines reviewed here offer similar functionality, with
537 particularly StaG-mwc having a strong overlap with nf-core/taxprofiler. Thus, users
538 in most cases will be able to select the pipeline depending on which framework they
539 feel most comfortable with. However the advantages of nf-core/taxprofiler mainly

²As demonstrated in this blogpost from Paweł Przytuła: <https://web.archive.org/web/20230320223436/https://appsilon.com/reproducible-research-when-your-results-cant-be-reproduced/> (Accessed 2023-08-25)

come from the offering of the greatest choice of tools, as well the particular benefits provided by Nextflow. It provides the greatest number of computational infrastructure types the pipeline can be executed on, and container systems can be used to ensure reproducibility, as well the support of the nf-core community due to the centralised pool of ‘plug-and-play’ modules to make it easier to update the pipeline over time to add new tools classifiers.

The functionality offered by other pipelines not currently supported by nf-core/taxprofiler include sequencing saturation estimation (StaG-mwc), taxonomy-free composition comparison (StaG-mwc), functional profiling (StaG-mwc), *de novo* assembly (sunbeam), and reference mapping (StaG-mwc, sunbeam). We do not plan to support *de novo* assembly or functional profiling in nf-core/taxprofiler as we feel these are already better served by other existing dedicated pipelines within the nf-core ecosystem: nf-core/mag for *de novo* assembly, (Krakau et al. 2022) and nf-core/funcscan for functional profiling (<https://nf-co.re/funcscan>), as well as elsewhere e.g. MetaWRAP (Uritskiy, DiRuggiero, and Taylor 2018).

References

- Andrews, Simon. 2010. “FastQC: A Quality Control Tool for High Throughput Sequence Data.” <http://www.bioinformatics.babraham.ac.uk/projects/fastqc/>.
- Beber, Moritz E, Maxime Borry, Sofia Stamouli, and James A Fellows Yates. 2023. “TAXPASTA: TAXonomic Profile Aggregation and STandardisation.” *Journal of Open Source Software* 8 (87): 5627. <https://doi.org/10.21105/joss.05627>.
- Bengtsson-Palme, Johan, Martin Hartmann, Karl Martin Eriksson, Chandan Pal, Kaisa Thorell, Dan Göran Joakim Larsson, and Rolf Henrik Nilsson. 2015. “METAXA2: Improved Identification and Taxonomic Classification of Small and Large Subunit rRNA in Metagenomic Data.” *Molecular Ecology Resources* 15 (6): 1403–14. <https://doi.org/10.1111/1755-0998.12399>.
- Blanco-Míguez, Aitor, Francesco Beghini, Fabio Cumbo, Lauren J McIver, Kelsey N Thompson, Moreno Zolfo, Paolo Manghi, et al. 2023. “Extending and Improving Metagenomic Taxonomic Profiling with Uncharacterized Species Using MetaPhlAn 4.” *Nature Biotechnology*, February, 1–12. <https://doi.org/10.1038/s41587-023-01688-w>.
- Bolger, Anthony M, Marc Lohse, and Bjoern Usadel. 2014. “Trimmomatic: A Flexible Trimmer for Illumina Sequence Data.” *Bioinformatics (Oxford, England)* 30 (15): 2114–20. <https://doi.org/10.1093/bioinformatics/btu170>.
- Boulund, Fredrik, Aron Arzoomand, Justine Debelius, chrsb, and Lisa Olsson. 2023. “Ctmbio/Stag-Mwc: StaG v0.7.0.” Zenodo. <https://doi.org/10.5281/ZENODO.8032462>.
- Breitwieser, F P, D N Baker, and S L Salzberg. 2018. “KrakenUniq: Confident and Fast Metagenomics Classification Using Unique k-Mer Counts.” *Genome Biology* 19 (1): 198. <https://doi.org/10.1186/s13059-018-1568-0>.
- Breitwieser, Florian P, Jennifer Lu, and Steven L Salzberg. 2019. “A Review of Methods and Databases for Metagenomic Classification and Assembly.” *Briefings in Bioinformatics* 20 (4): 1125–36. <https://doi.org/10.1093/bib/bbx120>.

583 Breitwieser, Florian P, Mihaela Pertea, Aleksey Zimin, and Steven L Salzberg. 2019.
584 “Human Contamination in Bacterial Genomes Has Created Thousands of Spurious
585 Proteins.” *Genome Research* 29 (May): 954–60. [https://doi.org/10.1101/gr.245373.](https://doi.org/10.1101/gr.245373.118)
586 118.

587 Buchfink, Benjamin, Klaus Reuter, and Hajk-Georg Drost. 2021. “Sensitive Protein
588 Alignments at Tree-of-Life Scale Using DIAMOND.” *Nature Methods* 18 (4): 366–
589 68. <https://doi.org/10.1038/s41592-021-01101-x>.

590 Bushnell, Brian. 2022. “BBMap.” <https://sourceforge.net/projects/bbmap/>.

591 Cantu, Vito Adrian, Jeffrey Sadural, and Robert Edwards. 2019. “PRINSEQ++, a Multi-
592 Threaded Tool for Fast and Efficient Quality Control and Preprocessing of Se-
593 quencing Datasets.” e27553v1. PeerJ Preprints; PeerJ Inc. [https://doi.org/10.7287/](https://doi.org/10.7287/peerj.preprints.27553v1)
594 [peerj.preprints.27553v1](https://doi.org/10.7287/peerj.preprints.27553v1).

595 Chen, Shifu, Yanqing Zhou, Yaru Chen, and Jia Gu. 2018. “Fastp: An Ultra-Fast All-
596 in-One FASTQ Preprocessor.” *Bioinformatics* 34 (17): i884–90. [https://doi.org/10.](https://doi.org/10.1093/bioinformatics/bty560)
597 [1093/bioinformatics/bty560](https://doi.org/10.1093/bioinformatics/bty560).

598 Chiu, Charles Y, and Steven A Miller. 2019. “Clinical Metagenomics.” *Nature Reviews.*
599 *Genetics* 20 (6): 341–55. <https://doi.org/10.1038/s41576-019-0113-7>.

600 Clarke, Erik L, Louis J Taylor, Chunyu Zhao, Andrew Connell, Jung-Jin Lee, Bryton
601 Fett, Frederic D Bushman, and Kyle Bittinger. 2019. “Sunbeam: An Extensible
602 Pipeline for Analyzing Metagenomic Sequencing Experiments.” *Microbiome* 7 (1):
603 46. <https://doi.org/10.1186/s40168-019-0658-x>.

604 Danecek, Petr, James K Bonfield, Jennifer Liddle, John Marshall, Valeriu Ohan, Mar-
605 tin O Pollard, Andrew Whitwham, et al. 2021. “Twelve Years of SAMtools and
606 BCftools.” *GigaScience* 10 (2). <https://doi.org/10.1093/gigascience/giab008>.

607 DeSantis, T Z, P Hugenholtz, N Larsen, M Rojas, E L Brodie, K Keller, T Huber, D
608 Dalevi, P Hu, and G L Andersen. 2006. “Greengenes, a Chimera-Checked 16S
609 rRNA Gene Database and Workbench Compatible with ARB.” *Applied and Envi-*
610 *ronmental Microbiology* 72 (7): 5069–72. <https://doi.org/10.1128/AEM.03006-05>.

611 Di Tommaso, Paolo, Maria Chatzou, Evan W Floden, Pablo Prieto Barja,
612 Emilio Palumbo, and Cedric Notredame. 2017. “Nextflow Enables Repro-
613 ducible Computational Workflows.” *Nature Biotechnology* 35 (4): 316–19.
614 <https://doi.org/10.1038/nbt.3820>.

615 Dilthey, Alexander T, Chirag Jain, Sergey Koren, and Adam M Phillippy. 2019. “Strain-
616 Level Metagenomic Assignment and Compositional Estimation for Long Reads
617 with MetaMaps.” *Nature Communications* 10 (1): 3066. [https://doi.org/10.1038/](https://doi.org/10.1038/s41467-019-10934-2)
618 [s41467-019-10934-2](https://doi.org/10.1038/s41467-019-10934-2).

619 Eloe-Fadros, Emiley A, Natalia N Ivanova, Tanja Woyke, and Nikos C Kyrpides. 2016.
620 “Metagenomics Uncovers Gaps in Amplicon-Based Detection of Microbial Diver-
621 sity.” *Nature Microbiology* 1 (4): 15032. <https://doi.org/10.1038/nmicrobiol.2015.32>.

622 Ewels, Philip A, Alexander Peltzer, Sven Fillinger, Harshil Patel, Johannes
623 Alneberg, Andreas Wilm, Maxime Ulysse Garcia, Paolo Di Tommaso,
624 and Sven Nahnsen. 2020. “The Nf-Core Framework for Community-
625 Curated Bioinformatics Pipelines.” *Nature Biotechnology* 38 (3): 276–78.
626 <https://doi.org/10.1038/s41587-020-0439-x>.

627 Govender, Kumeren N, and David W Eyre. 2022. “Benchmarking Taxonomic Classi-
628 fiers with Illumina and Nanopore Sequence Data for Clinical Metagenomic Diag-

629 nostic Applications.” *Microbial Genomics* 8 (10): 000886. [https://doi.org/10.1099/](https://doi.org/10.1099/mgen.0.000886)
630 [mgen.0.000886](https://doi.org/10.1099/mgen.0.000886).

631 Grüning, Björn, Ryan Dale, Andreas Sjödin, Brad A Chapman, Jillian Rowe,
632 Christopher H Tomkins-Tinch, Renan Valieris, Johannes Köster, and Bio-
633 conda Team. 2018. “Bioconda: Sustainable and Comprehensive Soft-
634 ware Distribution for the Life Sciences.” *Nature Methods* 15 (7): 475–76.
635 <https://doi.org/10.1038/s41592-018-0046-7>.

636 Hillmann, Benjamin, Gabriel A Al-Ghalith, Robin R Shields-Cutler, Qiyun Zhu, Daryl
637 M Gohl, Kenneth B Beckman, Rob Knight, and Dan Knights. 2018. “Evaluating the
638 Information Content of Shallow Shotgun Metagenomics.” *mSystems* 3 (6). <https://doi.org/10.1128/mSystems.00069-18>.

639 Kim, Daehwan, Li Song, Florian P Breitwieser, and Steven L Salzberg. 2016. “Cen-
640 trifuge: Rapid and Sensitive Classification of Metagenomic Sequences.” *Genome*
641 *Research* 26 (12): 1721–29. <https://doi.org/10.1101/gr.210641.116>.

642 Krakau, Sabrina, Daniel Straub, Hadrien Gourel, Gisela Gabernet, and Sven Nahnsen.
643 2022. “NF-Core/Mag: A Best-Practice Pipeline for Metagenome Hybrid Assembly
644 and Binning.” *NAR Genomics and Bioinformatics* 4 (1). [https://doi.org/10.1093/](https://doi.org/10.1093/nargab/lqac007)
645 [nargab/lqac007](https://doi.org/10.1093/nargab/lqac007).

646 Kryukov, Kirill, and Tadashi Imanishi. 2016. “Human Contamination in Public
647 Genome Assemblies.” *PloS One* 11 (9): e0162424. [https://doi.org/10.1371/journal.](https://doi.org/10.1371/journal.pone.0162424)
648 [pone.0162424](https://doi.org/10.1371/journal.pone.0162424).

649 Langmead, Ben, and Steven L Salzberg. 2012. “Fast Gapped-Read Alignment with
650 Bowtie 2.” *Nature Methods* 9 (4): 357–59. <https://doi.org/10.1038/nmeth.1923>.

651 Langmead, Ben, Christopher Wilks, Valentin Antonescu, and Rone Charles. 2019.
652 “Scaling Read Aligners to Hundreds of Threads on General-Purpose Processors.”
653 *Bioinformatics* 35 (3): 421–32. <https://doi.org/10.1093/bioinformatics/bty648>.

654 Li, Heng. 2018. “Minimap2: Pairwise Alignment for Nucleotide Sequences.” *Bioinfor-*
655 *matics* 34 (18): 3094–3100. <https://doi.org/10.1093/bioinformatics/bty191>.

656 Li, Heng, Bob Handsaker, Alec Wysoker, Tim Fennell, Jue Ruan, Nils Homer, Gabor
657 Marth, Goncalo Abecasis, Richard Durbin, and 1000 Genome Project Data Pro-
658 cessing Subgroup. 2009. “The Sequence Alignment/Map Format and SAMtools.”
659 *Bioinformatics* 25 (16): 2078–79. <https://doi.org/10.1093/bioinformatics/btp352>.

660 Longo, Mark S, Michael J O’Neill, and Rachel J O’Neill. 2011. “Abundant Human
661 DNA Contamination Identified in Non-Primate Genome Databases.” *PloS One* 6
662 (2): e16410. <https://doi.org/10.1371/journal.pone.0016410>.

663 Lu, Jennifer, Florian P Breitwieser, Peter Thielen, and Steven L Salzberg. 2017.
664 “Bracken: Estimating Species Abundance in Metagenomics Data.” *PeerJ*.
665 *Computer Science* 3 (e104): e104. <https://doi.org/10.7717/peerj-cs.104>.

666 Lynch, Michael D J, and Josh D Neufeld. 2015. “Ecology and Exploration of the Rare
667 Biosphere.” *Nature Reviews. Microbiology* 13 (4): 217–29. [https://doi.org/10.1038/](https://doi.org/10.1038/nrmicro3400)
668 [nrmicro3400](https://doi.org/10.1038/nrmicro3400).

669 McIntyre, Alexa B R, Rachid Ounit, Ebrahim Afshinnekoo, Robert J Prill, Elizabeth
670 Hénaff, Noah Alexander, Samuel S Minot, et al. 2017. “Comprehensive Bench-
671 marking and Ensemble Approaches for Metagenomic Classifiers.” *Genome Biology*
672 18 (1): 182. <https://doi.org/10.1186/s13059-017-1299-7>.

673 Menzel, Peter, Kim Lee Ng, and Anders Krogh. 2016. “Fast and Sensitive Taxonomic

675 Classification for Metagenomics with Kaiju.” *Nature Communications* 7 (April):
676 11257. <https://doi.org/10.1038/ncomms11257>.

677 Metwally, Ahmed A, Yang Dai, Patricia W Finn, and David L Perkins. 2016. “WEVOTE:
678 Weighted VOTing Taxonomic idEntification Method of Microbial Sequences.” *PloS*
679 *One* 11 (9): e0163527. <https://doi.org/10.1371/journal.pone.0163527>.

680 Meyer, Fernando, Adrian Fritz, Zhi-Luo Deng, David Koslicki, Till Robin Lesker,
681 Alexey Gurevich, Gary Robertson, et al. 2022. “Critical Assessment of
682 Metagenome Interpretation: The Second Round of Challenges.” *Nature Methods*
683 19 (4): 429–40. <https://doi.org/10.1038/s41592-022-01431-4>.

684 Mitchell, Alex L, Alexandre Almeida, Martin Beracochea, Miguel Boland, Josephine
685 Burgin, Guy Cochrane, Michael R Crusoe, et al. 2019. “MGnify: The Microbiome
686 Analysis Resource in 2020.” *Nucleic Acids Research*, November. <https://doi.org/10.1093/nar/gkz1035>.

687 Mölder, Felix, Kim Philipp Jablonski, Brice Letcher, Michael B Hall, Christopher H
688 Tomkins-Tinch, Vanessa Sochat, Jan Forster, et al. 2021. “Sustainable Data Anal-
689 ysis with Snakemake.” *F1000Research* 10 (January): 33. <https://doi.org/10.12688/f1000research.29032.2>.

690 Morais, Diego A A, João V F Cavalcante, Shênia S Monteiro, Matheus A B Pasquali,
691 and Rodrigo J S Dalmolin. 2022. “MEDUSA: A Pipeline for Sensitive Taxonomic
692 Classification and Flexible Functional Annotation of Metagenomic Shotgun Se-
693 quences.” *Frontiers in Genetics* 13 (March): 814437. <https://doi.org/10.3389/fgene.2022.814437>.

694 Nasko, Daniel J, Sergey Koren, Adam M Phillippy, and Todd J Treangen. 2018. “Ref-
695 Seq Database Growth Influences the Accuracy of k-Mer-Based Lowest Common
696 Ancestor Species Identification.” *Genome Biology* 19 (1): 165. <https://doi.org/10.1186/s13059-018-1554-6>.

697 Nayfach, Stephen, and Katherine S Pollard. 2016. “Toward Accurate and Quantitative
698 Comparative Metagenomics.” *Cell* 166 (5): 1103–16. <https://doi.org/10.1016/j.cell.2016.08.007>.

699 Ondov, Brian D, Nicholas H Bergman, and Adam M Phillippy. 2011. “Interactive
700 Metagenomic Visualization in a Web Browser.” *BMC Bioinformatics* 12 (1): 385.
701 <https://doi.org/10.1186/1471-2105-12-385>.

702 Piro, Vitor C, Temesgen H Dadi, Enrico Seiler, Knut Reinert, and Bernhard Y Renard.
703 2020. “Ganon: Precise Metagenomics Classification Against Large and up-to-Date
704 Sets of Reference Sequences.” *Bioinformatics (Oxford, England)* 36 (Suppl_1): i12–
705 20. <https://doi.org/10.1093/bioinformatics/btaa458>.

706 Piro, Vitor C, Marcel Matschkowski, and Bernhard Y Renard. 2017. “MetaMeta: Inte-
707 grating Metagenome Analysis Tools to Improve Taxonomic Profiling.” *Microbiome*
708 5 (1): 101. <https://doi.org/10.1186/s40168-017-0318-y>.

709 Pochon, Zoé, Nora Bergfeldt, Emrah Kırdök, Mário Vicente, Thijessen Naidoo, Tom
710 van der Valk, N Ezgi Altınışık, et al. 2022. “aMeta: An Accurate and Memory-
711 Efficient Ancient Metagenomic Profiling Workflow.” *bioRxiv*. <https://doi.org/10.1101/2022.10.03.510579>.

712 Portik, Daniel M, C Titus Brown, and N Tessa Pierce-Ward. 2022. “Evaluation of
713 Taxonomic Classification and Profiling Methods for Long-Read Shotgun Metage-
714 nomic Sequencing Datasets.” *BMC Bioinformatics* 23 (1): 541. <https://doi.org/10.1186/s12859-022-04541-1>.

1186/s12859-022-05103-0.

Quince, Christopher, Alan W Walker, Jared T Simpson, Nicholas J Loman, and Nicola Segata. 2017. "Shotgun Metagenomics, from Sampling to Analysis." *Nature Biotechnology* 35 (9): 833–44. <https://doi.org/10.1038/nbt.3935>.

Rodriguez-R, Luis M, Santosh Gunturu, James M Tiedje, James R Cole, and Konstantinos T Konstantinidis. 2018. "Nonpareil 3: Fast Estimation of Metagenomic Coverage and Sequence Diversity." *mSystems* 3 (3). <https://doi.org/10.1128/mSystems.00039-18>.

Rose, Rebecca, Olga Golosova, Dmitrii Sukhomlinov, Aleksey Tiunov, and Mattia Proserpi. 2019. "Flexible Design of Multiple Metagenomics Classification Pipelines with UGENE." *Bioinformatics (Oxford, England)* 35 (11): 1963–65. <https://doi.org/10.1093/bioinformatics/bty901>.

Ruscheweyh, Hans-Joachim, Alessio Milanese, Lucas Paoli, Nicolai Karcher, Quentin Clayssen, Marisa Isabell Keller, Jakob Wirbel, et al. 2022. "Cultivation-Independent Genomes Greatly Expand Taxonomic-Profiling Capabilities of mOTUs Across Various Environments." *Microbiome* 10 (1): 212. <https://doi.org/10.1186/s40168-022-01410-z>.

Schäffer, Alejandro A, Eric P Nawrocki, Yoon Choi, Paul A Kitts, Ilene Karsch-Mizrachi, and Richard McVeigh. 2018. "VecScreen_plus_taxonomy: Imposing a Tax(onomy) Increase on Vector Contamination Screening." *Bioinformatics (Oxford, England)* 34 (5): 755–59. <https://doi.org/10.1093/bioinformatics/btx669>.

Schloss, Patrick D, Sarah L Westcott, Thomas Ryabin, Justine R Hall, Martin Hartmann, Emily B Hollister, Ryan A Lesniewski, et al. 2009. "Introducing Mothur: Open-Source, Platform-Independent, Community-Supported Software for Describing and Comparing Microbial Communities." *Applied and Environmental Microbiology* 75 (23): 7537–41. <https://doi.org/10.1128/AEM.01541-09>.

Schmieder, Robert, and Robert Edwards. 2011. "Quality Control and Preprocessing of Metagenomic Datasets." *Bioinformatics (Oxford, England)* 27 (6): 863–64. <https://doi.org/10.1093/bioinformatics/btr026>.

Schubert, Mikkel, Stinus Lindgreen, and Ludovic Orlando. 2016. "AdapterRemoval v2: Rapid Adapter Trimming, Identification, and Read Merging." *BMC Research Notes* 9 (February): 88. <https://doi.org/10.1186/s13104-016-1900-2>.

Sczyrba, Alexander, Peter Hofmann, Peter Belmann, David Koslicki, Stefan Janssen, Johannes Dröge, Ivan Gregor, et al. 2017. "Critical Assessment of Metagenome Interpretation-a Benchmark of Metagenomics Software." *Nature Methods* 14 (11): 1063–71. <https://doi.org/10.1038/nmeth.4458>.

Sena Brandine, Guilherme de, and Andrew D Smith. 2021. "Falco: High-Speed FastQC Emulation for Quality Control of Sequencing Data." *F1000Research* 8 (1874): 1874. <https://doi.org/10.12688/f1000research.21142.2>.

Sharpton, Thomas J. 2014. "An Introduction to the Analysis of Shotgun Metagenomic Data." *Frontiers in Plant Science* 5 (June): 209. <https://doi.org/10.3389/fpls.2014.00209>.

Shen, Wei, Hongyan Xiang, Tianquan Huang, Hui Tang, Mingli Peng, Dachuan Cai, Peng Hu, and Hong Ren. 2023. "KMCP: Accurate Metagenomic Profiling of Both Prokaryotic and Viral Populations by Pseudo-Mapping." *Bioinformatics* 39 (1): btac845. <https://doi.org/10.1093/bioinformatics/btac845>.

- 767 Sim, Mikang, Jongin Lee, Daehwan Lee, Daehong Kwon, and Jaebum Kim. 2020.
768 “TAMA: Improved Metagenomic Sequence Classification Through Meta-Analysis.”
769 *BMC Bioinformatics* 21 (1): 185. <https://doi.org/10.1186/s12859-020-3533-7>.
- 770 Straub, Daniel, Nia Blackwell, Adrian Langarica-Fuentes, Alexander Peltzer, Sven
771 Nahnsen, and Sara Kleindienst. 2020. “Interpretations of Environmental Micro-
772 bial Community Studies Are Biased by the Selected 16S rRNA (Gene) Amplicon
773 Sequencing Pipeline.” *Frontiers in Microbiology* 11 (October): 550420. <https://doi.org/10.3389/fmicb.2020.550420>.
- 774 Sun, Zheng, Shi Huang, Meng Zhang, Qiyun Zhu, Niina Haiminen, Anna Paola Car-
775 rieri, Yoshiki Vázquez-Baeza, et al. 2021. “Challenges in Benchmarking Metage-
776 nomic Profilers.” *Nature Methods* 18 (6): 618–26. <https://doi.org/10.1038/s41592-021-01141-3>.
- 777 Titus Brown, C, and Luiz Irber. 2016. “Sourmash: A Library for MinHash Sketching
778 of DNA.” *Journal of Open Source Software* 1 (5): 27. <https://doi.org/10.21105/joss.00027>.
- 779 Uritskiy, Gherman V, Jocelyne DiRuggiero, and James Taylor. 2018. “MetaWRAP-a
780 Flexible Pipeline for Genome-Resolved Metagenomic Data Analysis.” *Microbiome*
781 6 (1): 158. <https://doi.org/10.1186/s40168-018-0541-1>.
- 782 Vågøene, Åshild J, Alexander Herbig, Michael G Campana, Nelly M Robles García,
783 Christina Warinner, Susanna Sabin, Maria A Spyrou, et al. 2018. “Salmonella
784 Enterica Genomes from Victims of a Major Sixteenth-Century Epidemic in Mex-
785 ico.” *Nature Ecology & Evolution* 2 (3): 520–28. <https://doi.org/10.1038/s41559-017-0446-6>.
- 786 Veiga Leprevost, Felipe da, Björn A Gruning, Saulo Alves Aflitos, Hannes L
787 Röst, Julian Uszkoreit, Harald Barsnes, Marc Vaudel, et al. 2017. “Bio-
788 Containers: An Open-Source and Community-Driven Framework for Soft-
789 ware Standardization.” *Bioinformatics (Oxford, England)* 33 (16): 2580–82.
790 <https://doi.org/10.1093/bioinformatics/btx192>.
- 791 Wick, Ryan R, Louise M Judd, Claire L Gorrie, and Kathryn E Holt. 2017. “Comple-
792 ting Bacterial Genome Assemblies with Multiplex MinION Sequencing.” *Microbial*
793 *Genomics* 3 (10): e000132. <https://doi.org/10.1099/mgen.0.000132>.
- 794 Wood, Derrick E, Jennifer Lu, and Ben Langmead. 2019. “Improved Metagenomic
795 Analysis with Kraken 2.” *Genome Biology* 20 (1): 257. <https://doi.org/10.1186/s13059-019-1891-0>.
- 796 Wratten, Laura, Andreas Wilm, and Jonathan Göke. 2021. “Reproducible, Scalable,
797 and Shareable Analysis Pipelines with Bioinformatics Workflow Managers.” *Nature*
798 *Methods* 18 (10): 1161–68. <https://doi.org/10.1038/s41592-021-01254-9>.
- 799 Wright, Robyn J, André M Comeau, and Morgan G I Langille. 2023. “From Defaults to
800 Databases: Parameter and Database Choice Dramatically Impact the Performance
801 of Metagenomic Taxonomic Classification Tools.” *Microbial Genomics* 9 (3). <https://doi.org/10.1099/mgen.0.000949>.
- 802 Ye, Simon H, Katherine J Siddle, Daniel J Park, and Pardis C Sabeti. 2019. “Bench-
803 marking Metagenomics Tools for Taxonomic Classification.” *Cell* 178 (4): 779–94.
804 <https://doi.org/10.1016/j.cell.2019.07.010>.
- 805 Yilmaz, Pelin, Laura Wegener Parfrey, Pablo Yarza, Jan Gerken, Elmar Pruesse, Chris-
806 tian Quast, Timmy Schweer, Jörg Peplies, Wolfgang Ludwig, and Frank Oliver
807 Oliver

813 Glöckner. 2014. “The SILVA and ‘All-Species Living Tree Project (LTP)’ Taxo-
814 nomic Frameworks.” *Nucleic Acids Research* 42 (Database issue): D643–8. [https:](https://doi.org/10.1093/nar/gkt1209)
815 [//doi.org/10.1093/nar/gkt1209](https://doi.org/10.1093/nar/gkt1209).